

# Deep Learning Using



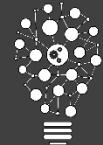
Masoud Pourreza



Pourreza.masoud@gmail.com



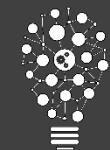
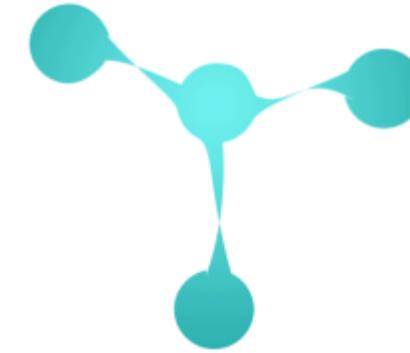
@masoudpz



HAMIM®

# Torch

- Open source machine learning library
- Scientific computing framework
- Based on the Lua !!
- Research Companies and Labs
- Facebook AI research, Google + Deep mind , Twitter, NVIDIA
- Started at 2002
- [website](#)



HAMIM®

# PyTorch

- PyTorch is a python package that provides two high-level features:
  - Tensor computation (like numpy) with strong GPU acceleration
  - Deep Neural Networks built on a tape-based autograd system
- PyTorch provides Tensors that can live either on the CPU or the GPU
- PyTorch is a python **library**
- More Pythonic (imperative)
- Flexible
- Intuitive and cleaner code
- Easy to debug
- More Neural Networkic
- Write code as the network works
- forward/backward

# PyTorch packages

Package	Description
torch	a Tensor library like NumPy, with strong GPU support
torch.autograd	a tape based automatic differentiation library that supports all differentiable Tensor operations in torch
torch.nn	a neural networks library deeply integrated with autograd designed for maximum flexibility
torch.optim	an optimization package to be used with torch.nn with standard optimization methods such as SGD, RMSProp, LBFGS, Adam etc.
torch.multiprocessing	python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and hogwild training.
torch.utils	DataLoader, Trainer and other utility functions for convenience

# Installation requirements

[Python on ubuntu](#)

[Pip on ubuntu](#)

[conda on ubuntu](#)

[Compile pytorch from source](#)

[Python on windows](#)

[Pip on windows](#)

[conda on windows](#)



HAMIM®

# PyTorch Installation

<http://pytorch.org>

## Get Started.

Select your preferences, then run the PyTorch install command.

Please ensure that you are on the latest pip and numpy packages.

Anaconda is our recommended package manager



Run this command:

```
pip3 install http://download.pytorch.org/whl/cu80/torch-0.3.0.post4-cp35-cp35m-linux_x86_64.whl  
pip3 install torchvision
```

[Click here for previous versions of PyTorch](#)

[Installation on windows using codna](#)

# Linear Model Example

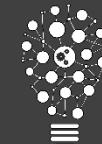
What would be the grade if I study 4 hours?

Hours (x)	Points (y)
1	2
2	4
3	6
4	?

Supervised Learning

Training data

Test data



HAMIM®

# Linear Model

- What would be the best model for the data? Linear?

Hours (x)	Points (y)
1	2
2	4
3	6
4	?

$$\hat{y} = x * w + b$$

$x$       Linear       $\hat{y}$

# Training Loss (error)

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

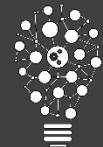
Hours, x	Points, y	Prediction, $y^w(w=3)$	Loss ( $w=3$ )
1	2	3	1
2	4	6	4
3	6	9	9
			mean=14/3



# Training Loss (error)

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

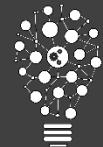
Hours, x	Points, y	Prediction, $y^{\wedge}(w=4)$	Loss (w=4)
1	2	4	4
2	4	8	16
3	6	12	36
		mean=56/3	



# Training Loss (error)

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

Hours, x	Points, y	Prediction, $y^{\wedge}(w=0)$	Loss (w=0)
1	2	0	4
2	4	0	16
3	6	0	36
			mean=56/3



# Training Loss (error)

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

Hours, x	Points, y	Prediction, $y^{\wedge}(w=1)$	Loss (w=1)
1	2	1	1
2	4	2	4
3	6	3	9
			mean=14/3



# Training Loss (error)

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

Hours, x	Points, y	Prediction, $y^{\wedge}(w=2)$	Loss (w=2)
1	2	2	0
2	4	4	0
3	6	6	0
			mean=0



HAMIM®

# Training Loss (error)

MSE mean square error

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

$$loss = \frac{1}{N} \sum_{n=1}^N (\hat{y}_n - y_n)^2$$

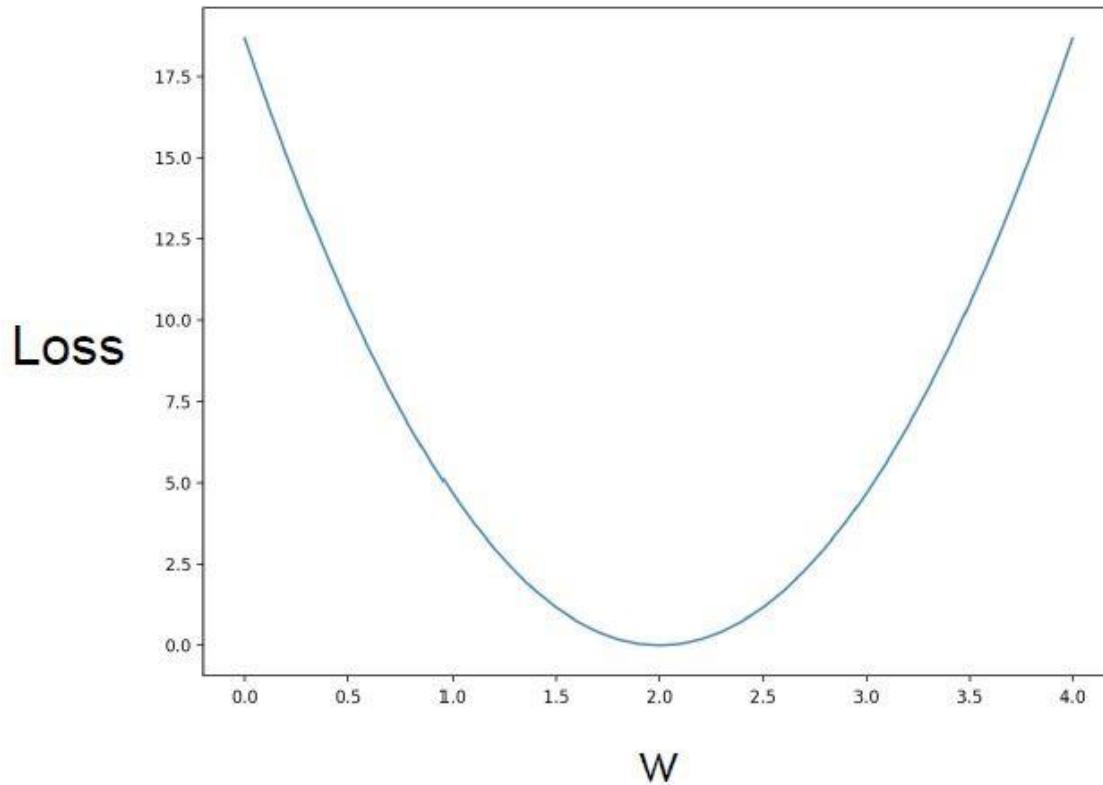
Hours, x	Loss (w=0)	Loss (w=1)	Loss (w=2)	Loss (w=3)	Loss (w=4)
1	4	1	0	1	4
2	16	4	0	4	16
3	36	9	0	9	36
	MSE=56/3=18.7	MSE=14/3=4.7	MSE=0	MSE=14/3=4.7	MSE=56/3=18.7



# Loss Graph

$$loss = \frac{1}{N} \sum_{n=1}^N (\hat{y}_n - y_n)^2$$

Loss (w=0)	Loss (w=1)	Loss (w=2)	Loss (w=3)	Loss (w=4)
mean=56/3=18.7	mean=14/3=4.7	mean=0	mean=14/3=4.7	mean=56/3=18.7



# Model & Loss

$$\hat{y} = x * w$$

$$loss = (\hat{y} - y)^2$$

```
w = 1.0 # a random guess: random value  
  
# our model for the forward pass  
def forward(x):  
    return x * w
```

```
# Loss function  
def loss(x, y):  
    y_pred = forward(x)  
    return (y_pred - y) * (y_pred - y)
```



# Compute loss for W

```
for w in np.arange(0.0, 4.1, 0.1):
    print("w=", w)
    l_sum = 0
    for x_val, y_val in zip(x_data, y_data):
        y_pred_val = forward(x_val)
        l = loss(x_val, y_val)
        l_sum += l
        print("\t", x_val, y_val, y_pred_val, l)

    print("MSE=", l_sum / 3)
```

```
w= 0.0
    1.0 2.0 0.0 4.0
    2.0 4.0 0.0 16.0
    3.0 6.0 0.0 36.0
NSE= 18.6666666667
w= 0.1
    1.0 2.0 0.1 3.61
    2.0 4.0 0.2 14.44
    3.0 6.0 0.3 32.49
NSE= 16.8466666667
w= 0.2
    1.0 2.0 0.2 3.24
    2.0 4.0 0.4 12.96
    3.0 6.0 0.6 29.16
NSE= 15.12
w= 0.3
    1.0 2.0 0.3 2.89
    2.0 4.0 0.6 11.56
    3.0 6.0 0.9 26.01
NSE= 13.4866666667
w= 0.4
    1.0 2.0 0.4 2.56
    2.0 4.0 0.8 10.24
    3.0 6.0 1.2 23.04
NSE= 11.9466666667
w= 0.5
    1.0 2.0 0.5 2.25
    2.0 4.0 1.0 9.0
    3.0 6.0 1.5 20.25
NSE= 10.5
```

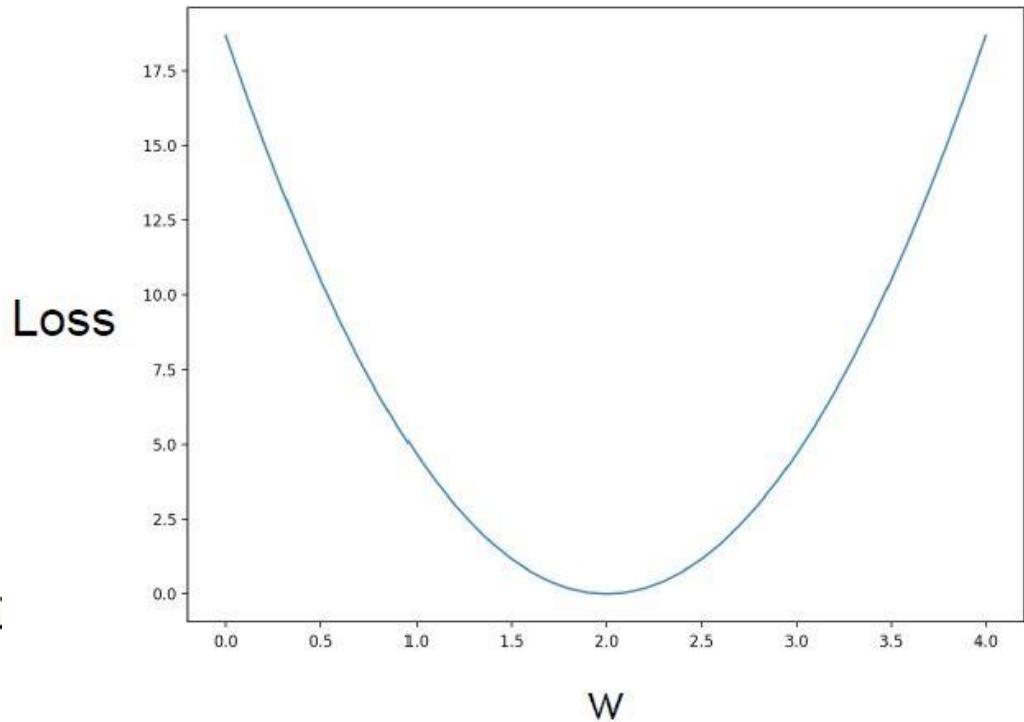


# Plot Graph

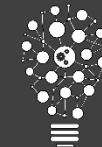
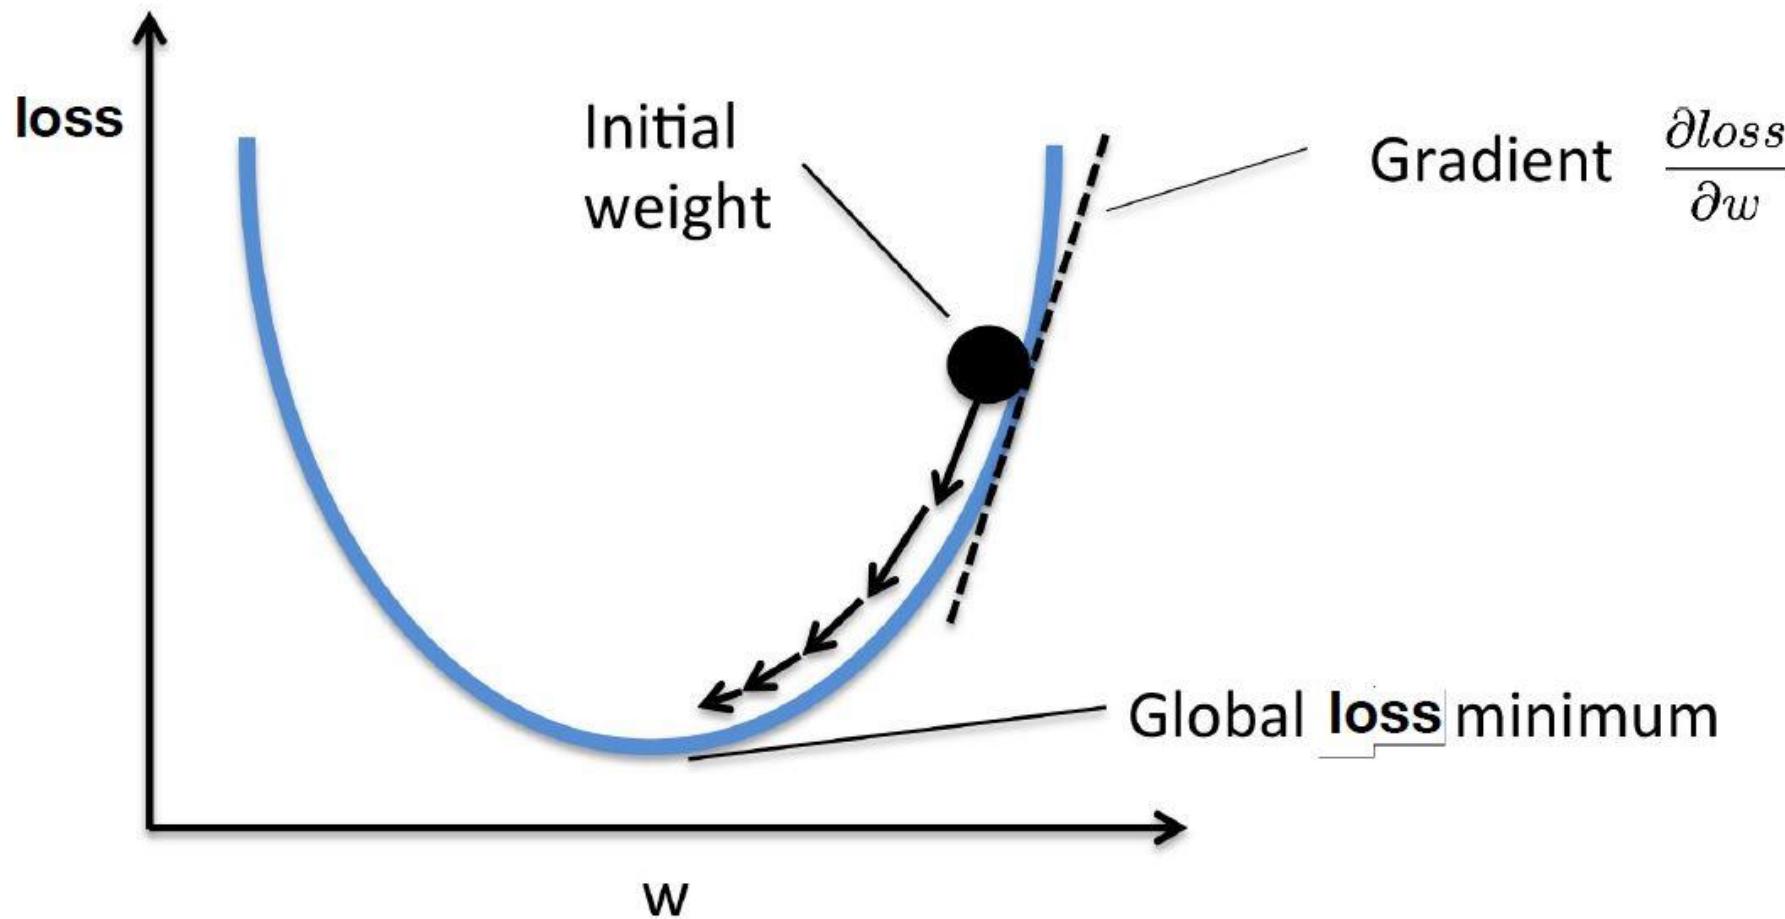
```
w_list = []
mse_list = []
for w in np.arange(0.0, 4.1, 0.1):
    print("w=", w)
    l_sum = 0
    for x_val, y_val in zip(x_data, y_data):
        y_pred_val = forward(x_val)
        l = loss(x_val, y_val)
        l_sum += l
        print("\t", x_val, y_val, y_pred_val, l)

    print("MSE=", l_sum / 3)
    w_list.append(w)
    mse_list.append(l_sum / 3)

plt.plot(w_list, mse_list)
plt.ylabel('Loss')
plt.xlabel('w')
plt.show()
```

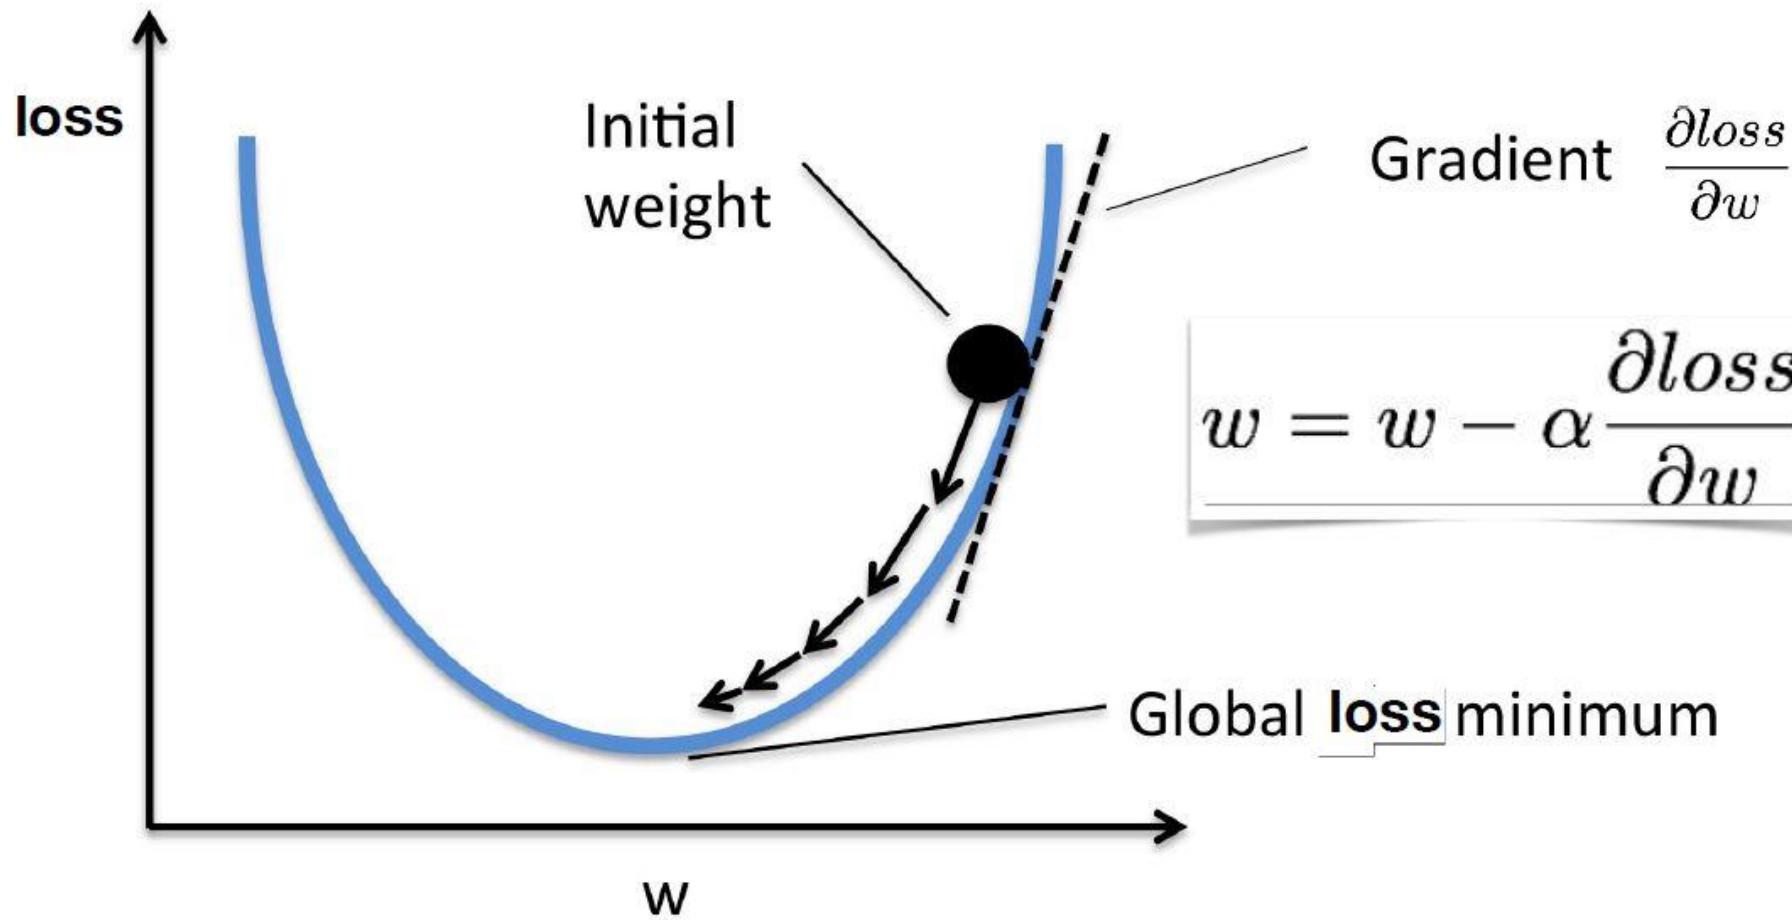


# Gradient Descent



HAMIM®

# Gradient Descent



# Derivative

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

$$w = w - \alpha \frac{\partial loss}{\partial w}$$

$$\frac{\partial loss}{\partial w} = ?$$

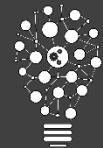


HAMIM®

# Derivative

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

$$\frac{\partial loss}{\partial w} = ?$$



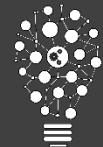
HAMIM®

# Derivative

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

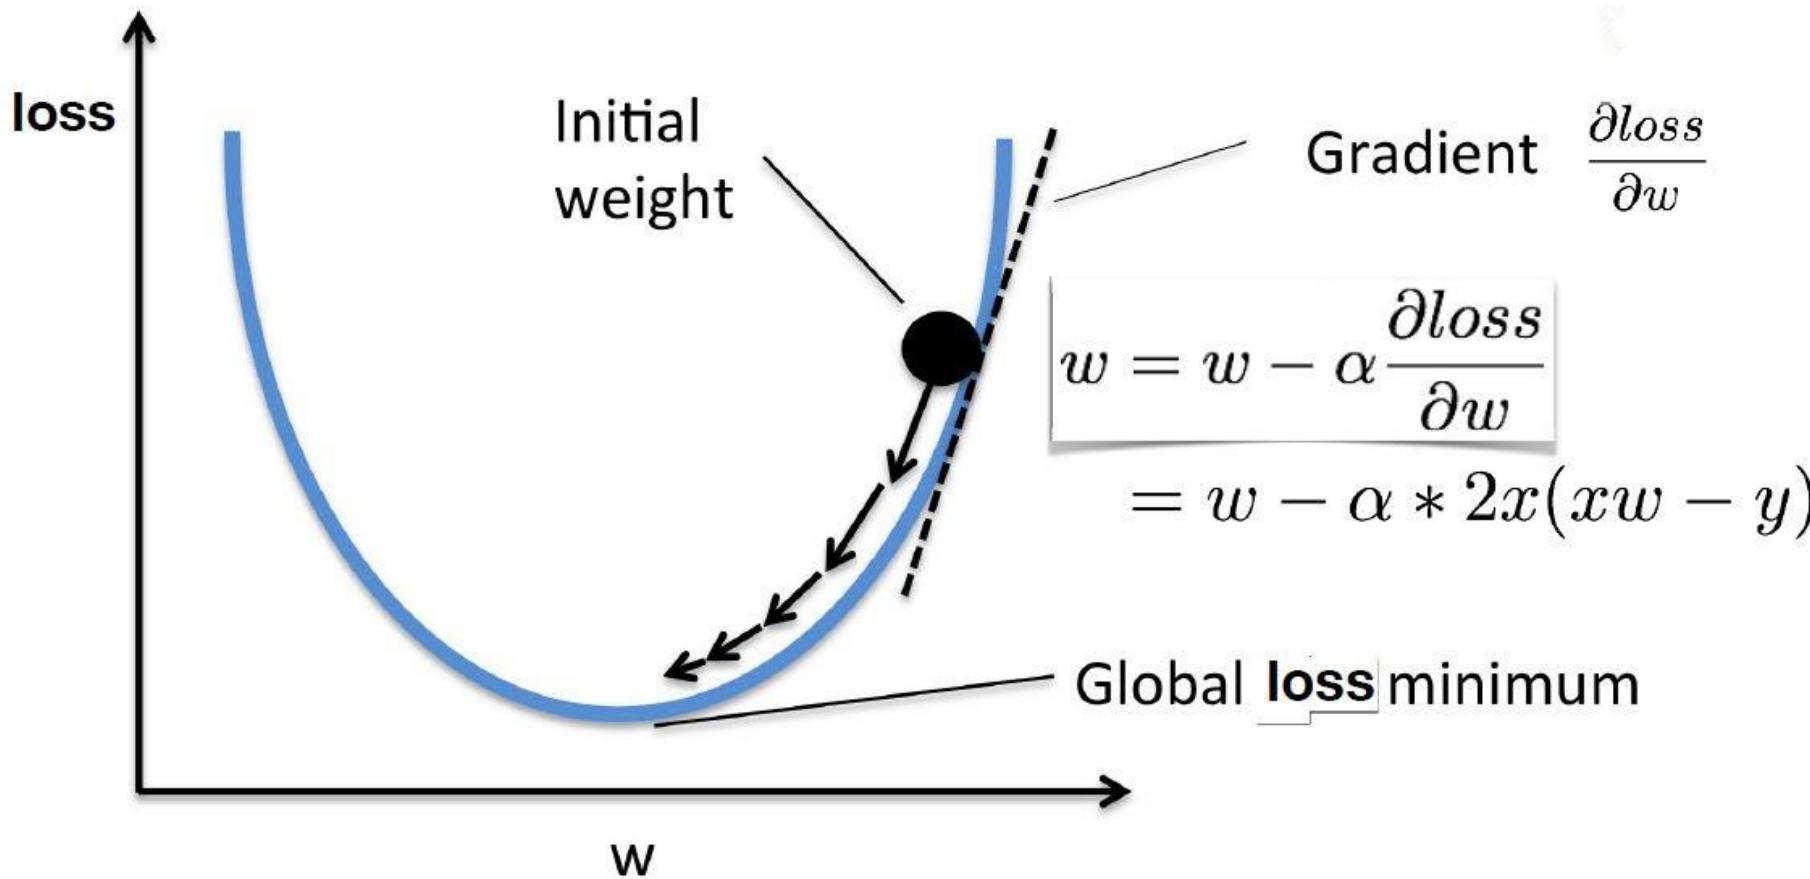
$$\frac{\partial loss}{\partial w} = ? \quad (U^2)' = 2 \times U \times U'$$

$$\frac{\partial loss}{\partial w} = 2 (x * w - y) * (x)$$



HAMIM®

# Gradient Descent



HAMIM®

# Data, Model , Loss and Gradient

```
x_data = [1.0, 2.0, 3.0]
y_data = [2.0, 4.0, 6.0]

w = 1.0 # a random guess: random value

# our model forward pass
def forward(x):
    return x * w

# Loss function
def loss(x, y):
    y_pred = forward(x)
    return (y_pred - y) * (y_pred - y)

# compute gradient
def gradient(x, y): # d_Loss/d_w
    return 2 * x * (x * w - y)
```



HAMIM®

# Refresh W

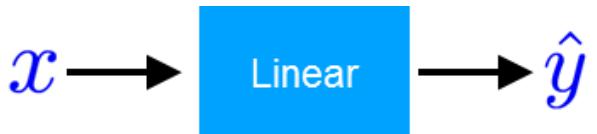
```
# Before training
print("predict (before training)", 4, forward(4))

# Training Loop
for epoch in range(100):
    for x_val, y_val in zip(x_data, y_data):
        grad = gradient(x_val, y_val)
        w = w - 0.01 * grad
        print("\tgrad: ", x_val, y_val, grad)
        l = loss(x_val, y_val)

    print("progress:", epoch, "w=", w, "loss=", l)

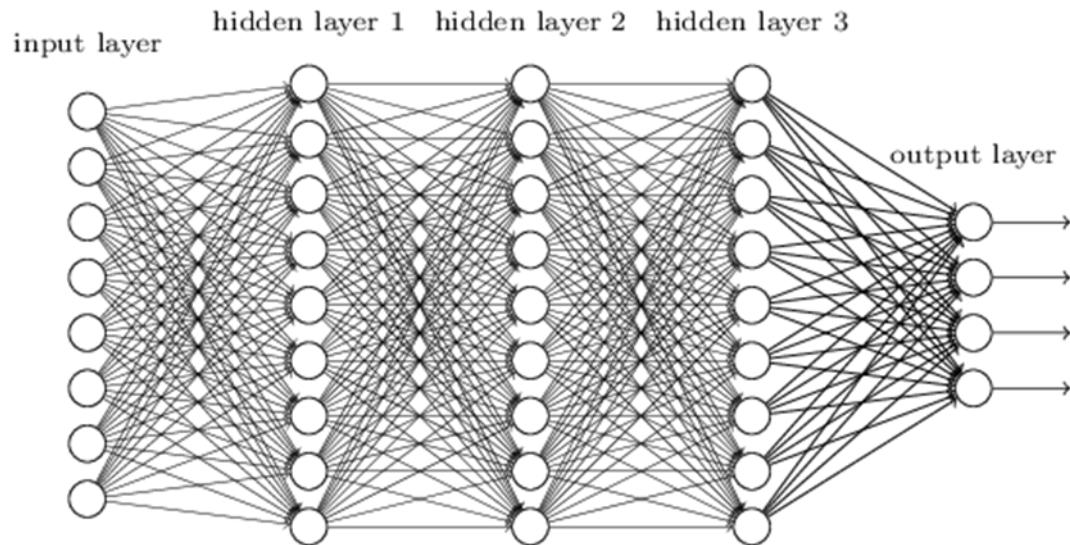
# After training
print("predict (after training)", "4 hours", forward(4))
```

# Back-propagation & Autograd



Gradient of **loss**  
with respect to  $w$   $\frac{\partial \text{loss}}{\partial w} = ?$

```
# compute gradient
def gradient(x, y): # d_Loss/d_w
    return 2 * x * (x * w - y)
```



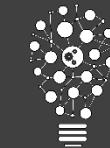
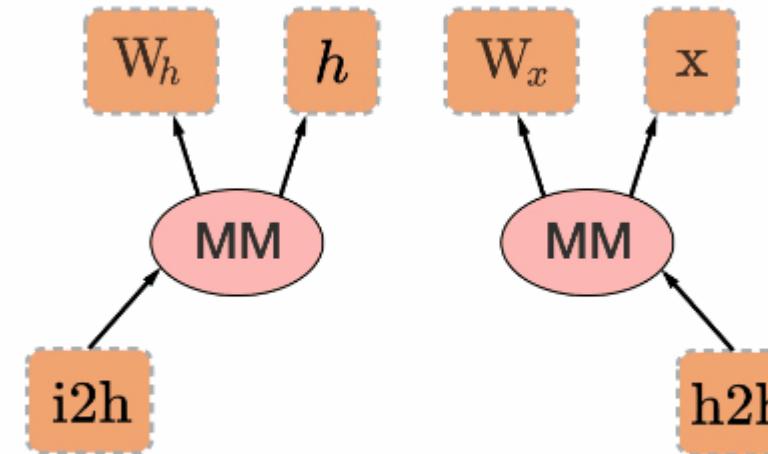
Gradient of **loss**  
with respect to  $w$   $\frac{\partial \text{loss}}{\partial w} = ?$

# Method? Computational Graph + Chain Rule



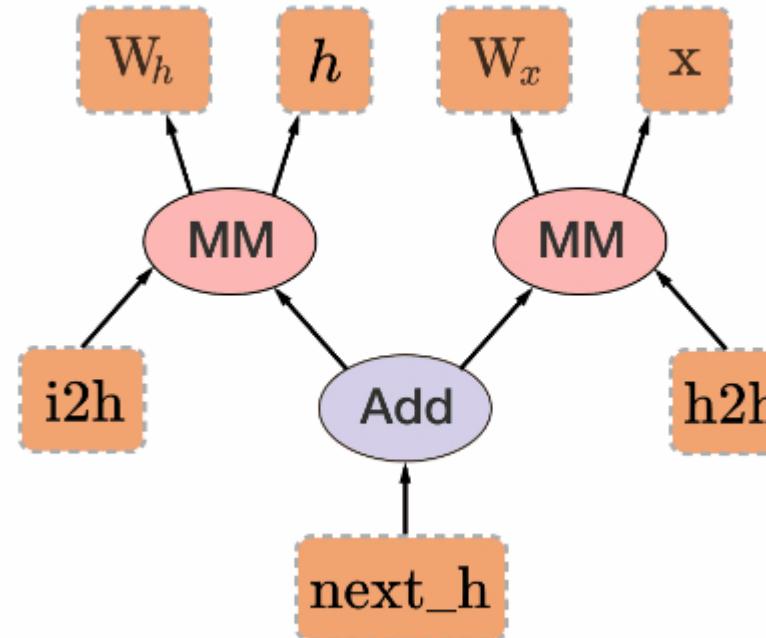
HAMIM®

# Method? Computational Graph + Chain Rule



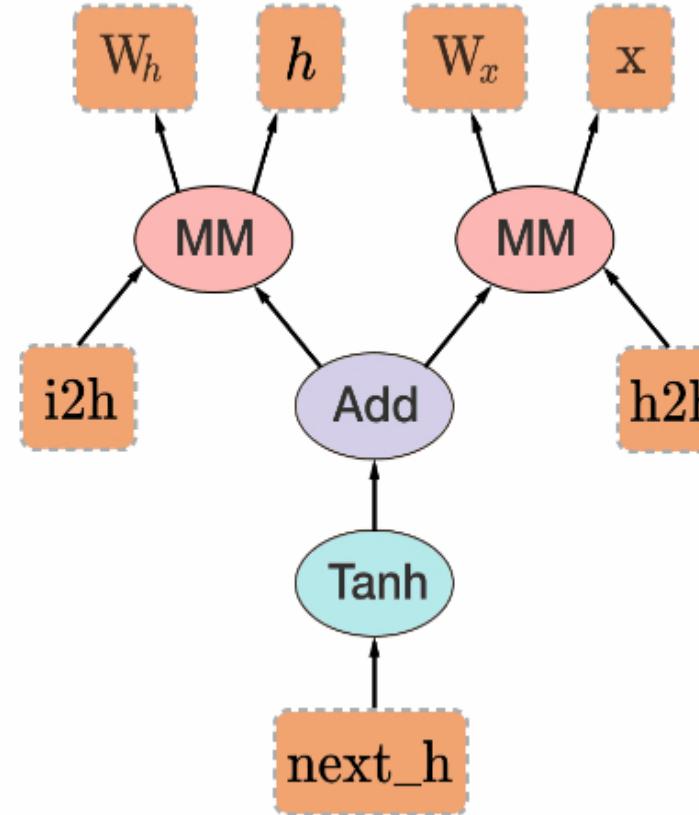
HAMIM®

# Method? Computational Graph + Chain Rule



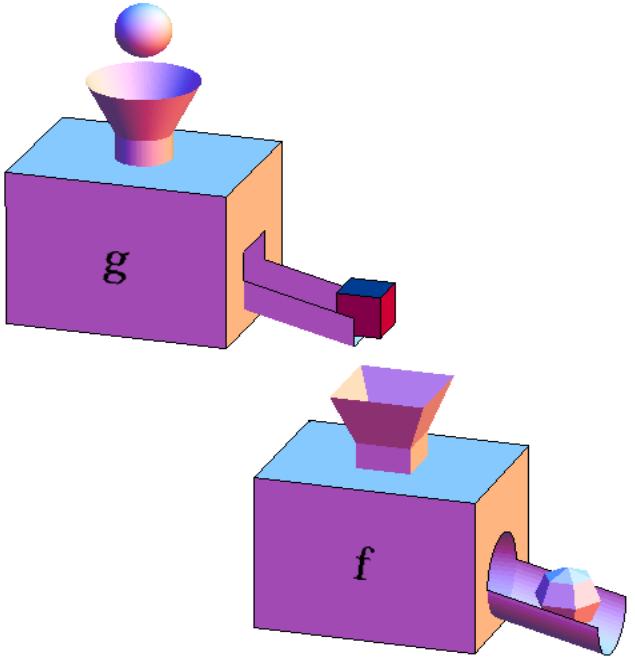
HAMIM®

# Method? Computational Graph + Chain Rule



HAMIM®

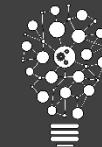
# Chain Rule



## The Chain Rule

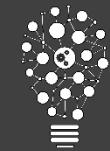
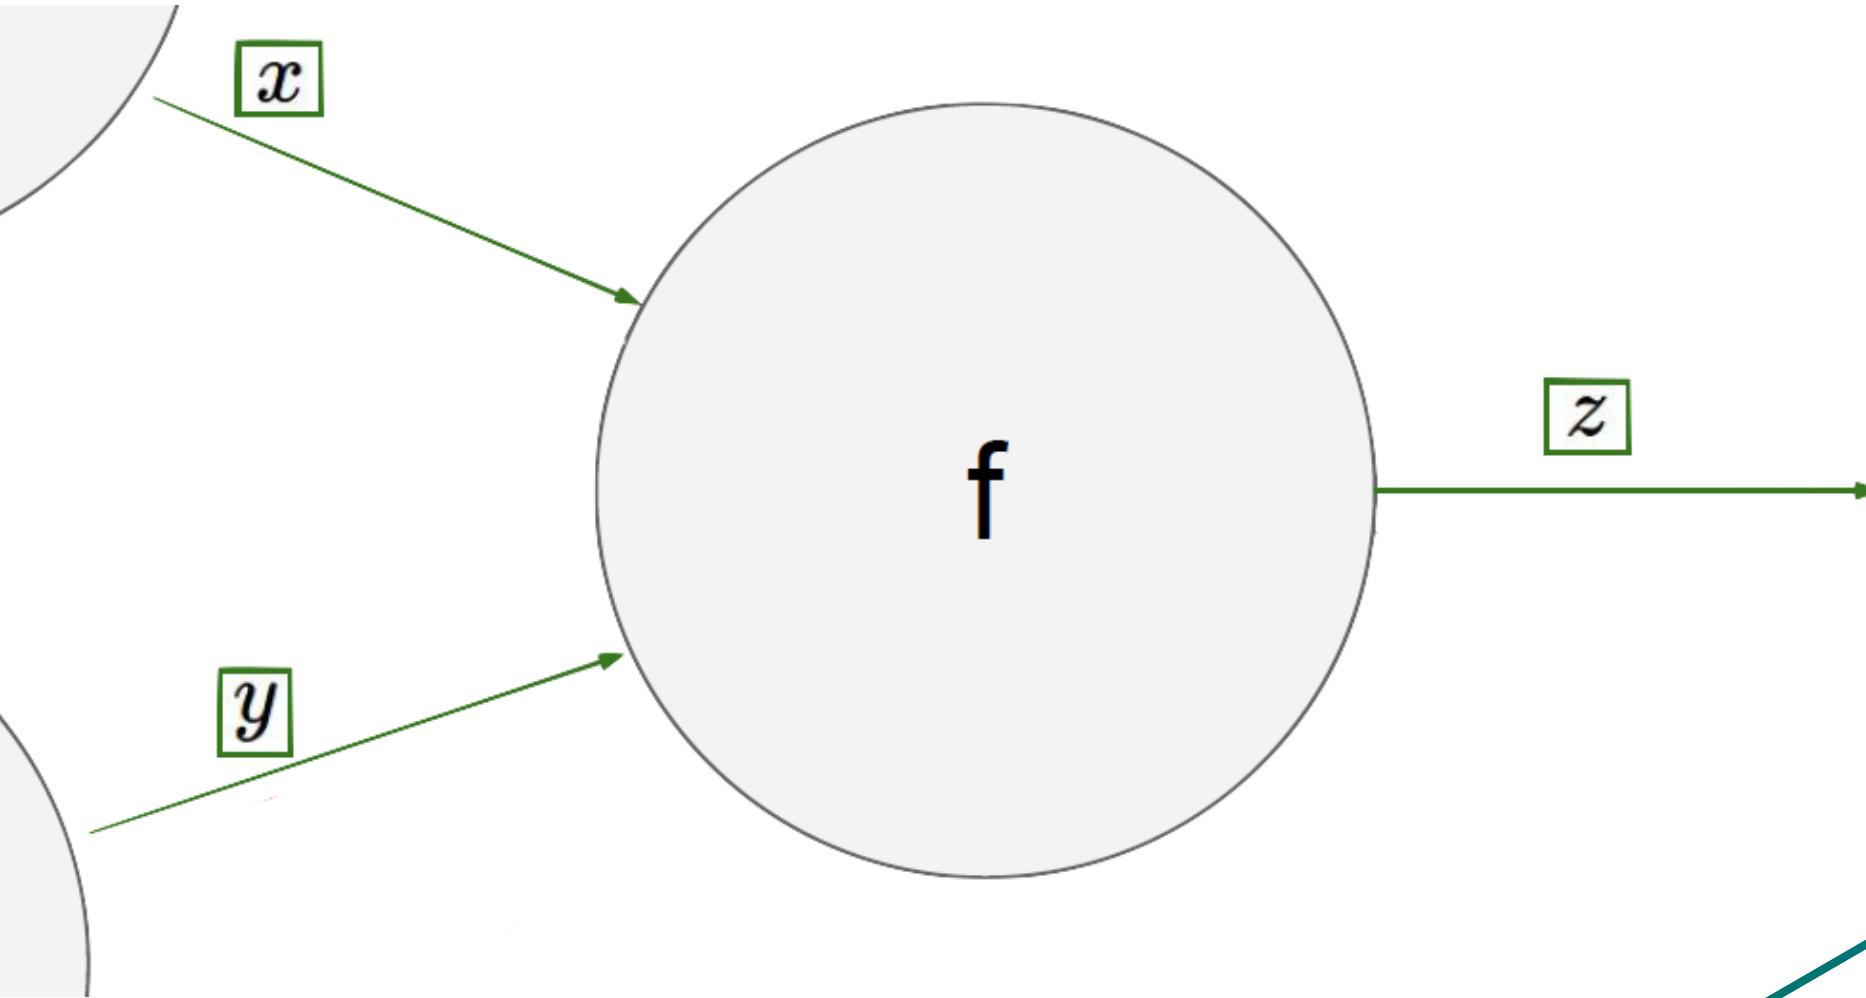
$$f = f(g) ; \quad g = g(x)$$

$$\frac{df}{dx} = \frac{df}{dg} \frac{dg}{dx}$$



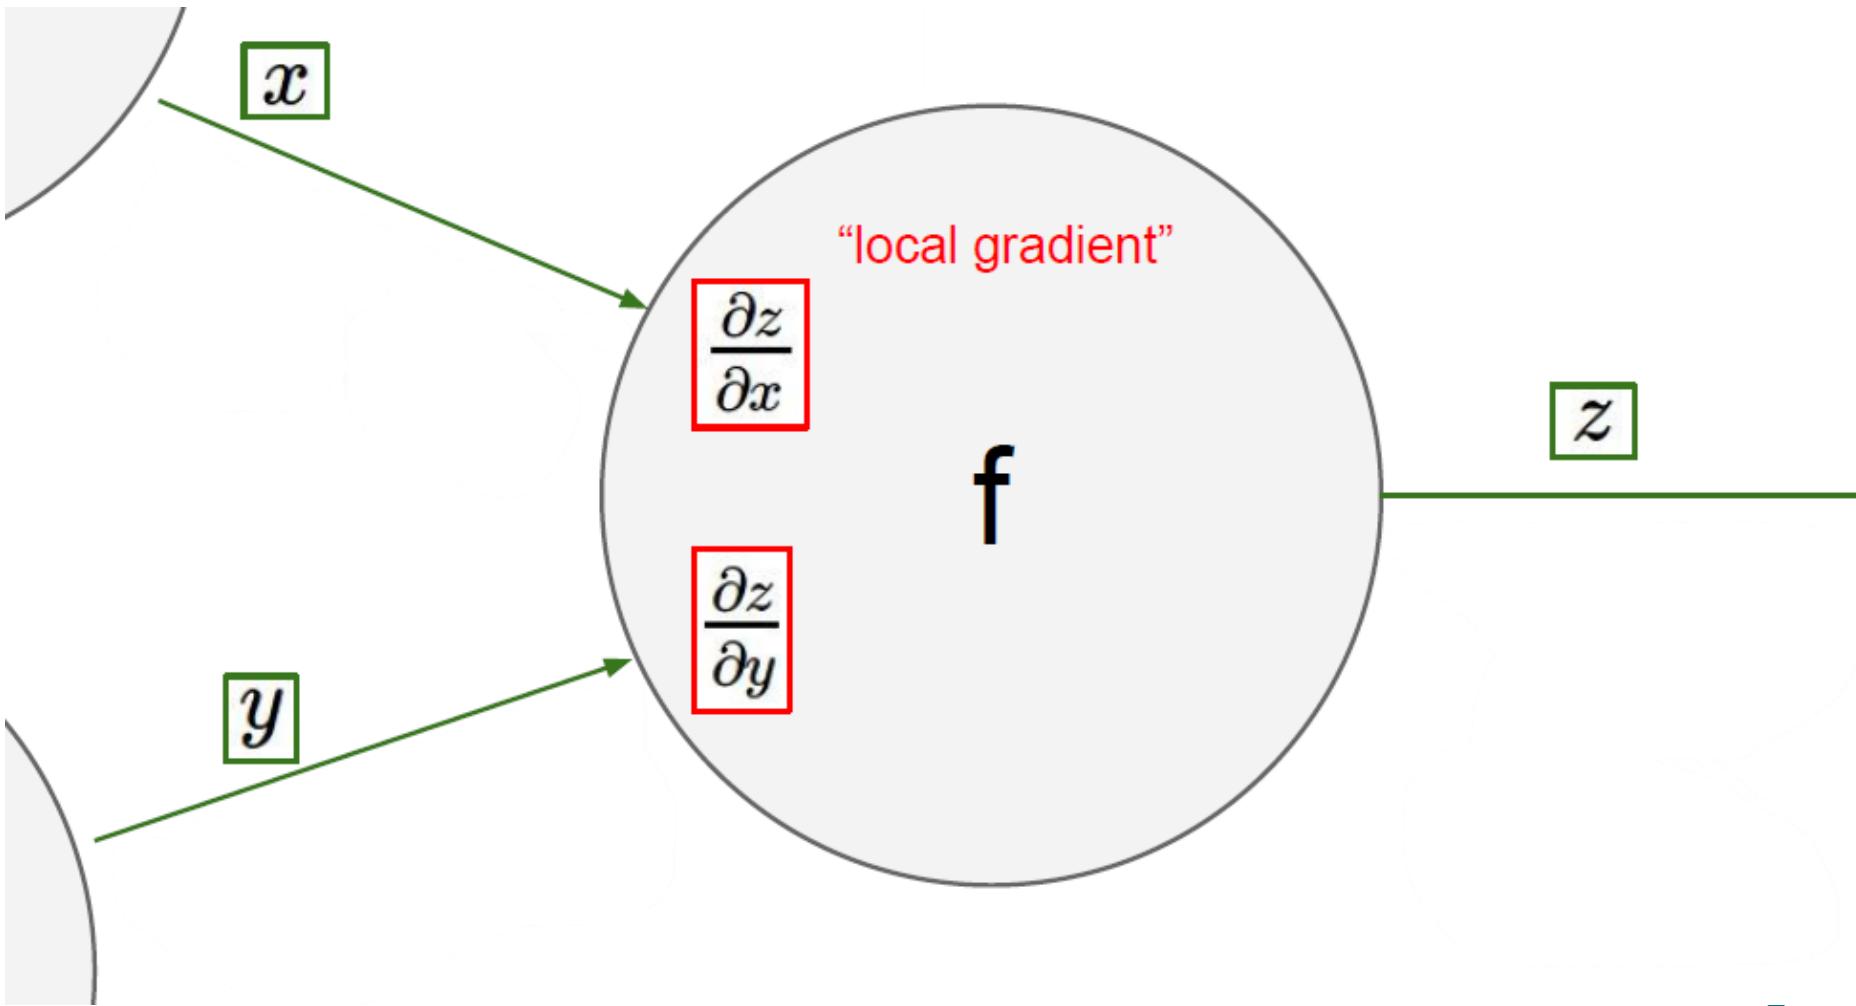
HAMIM®

# Chain Rule



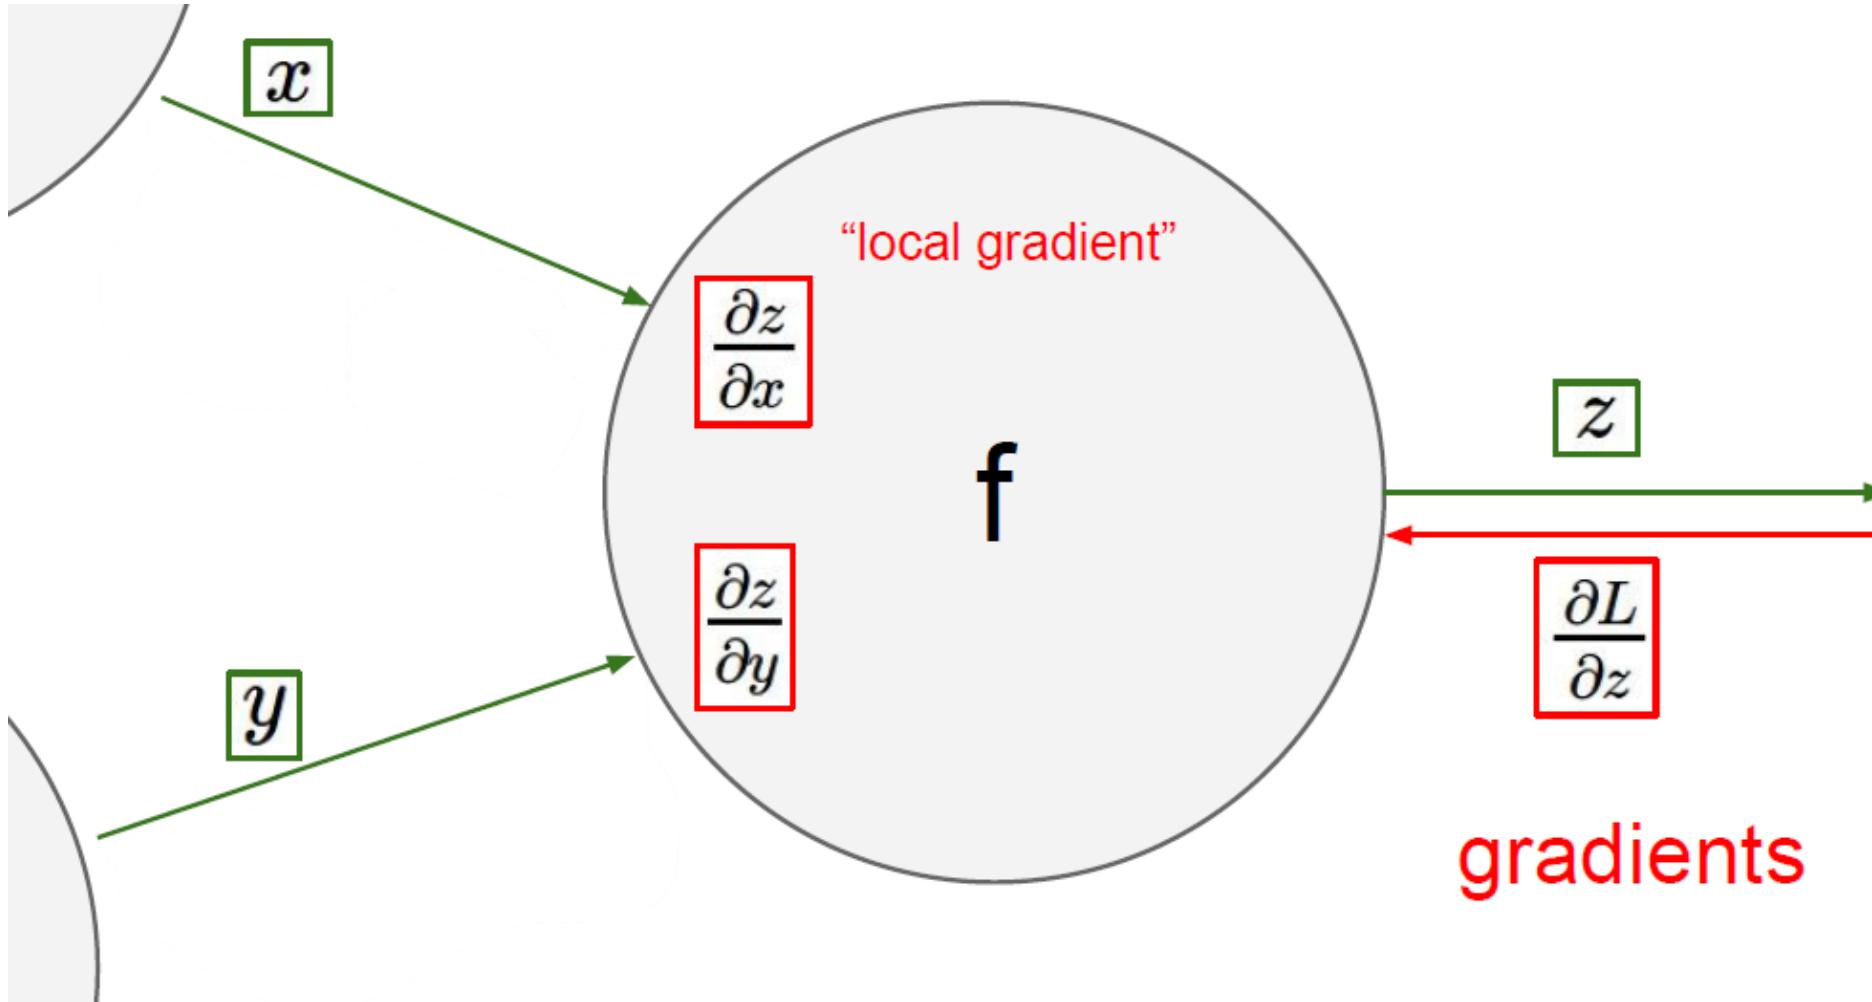
HAMIM®

# Chain Rule

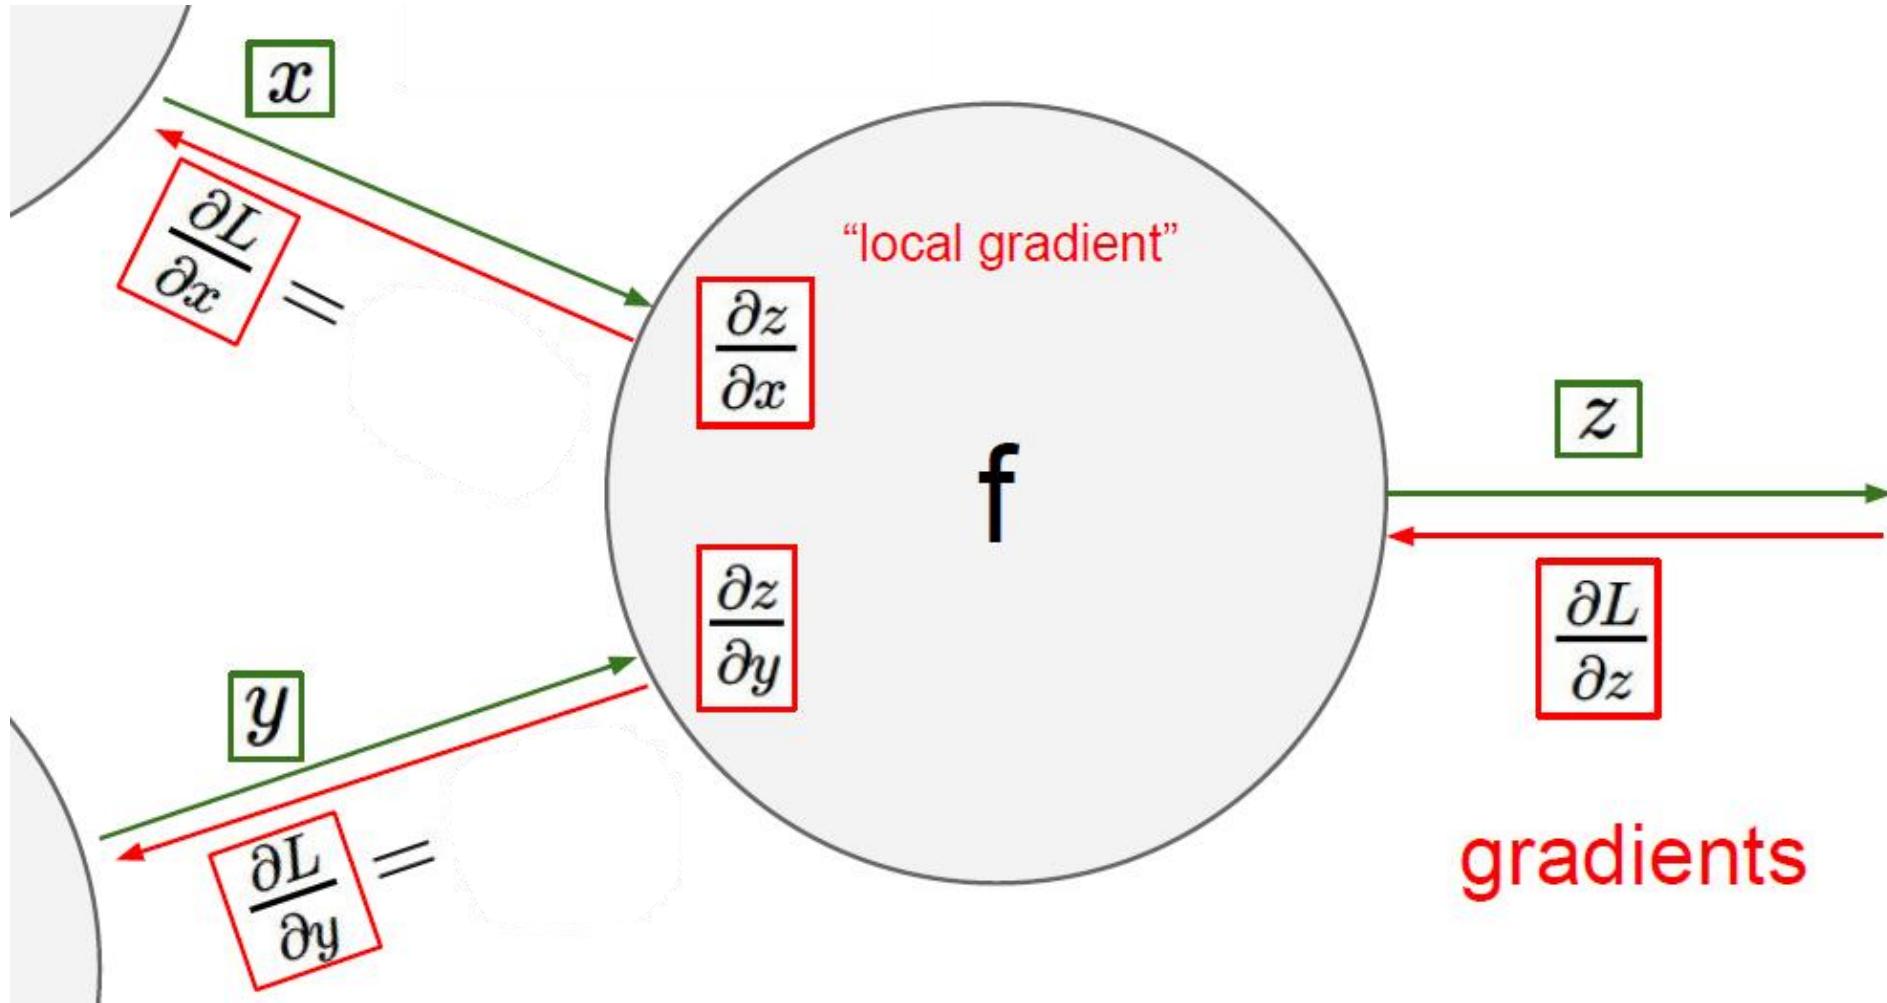


HAMIM®

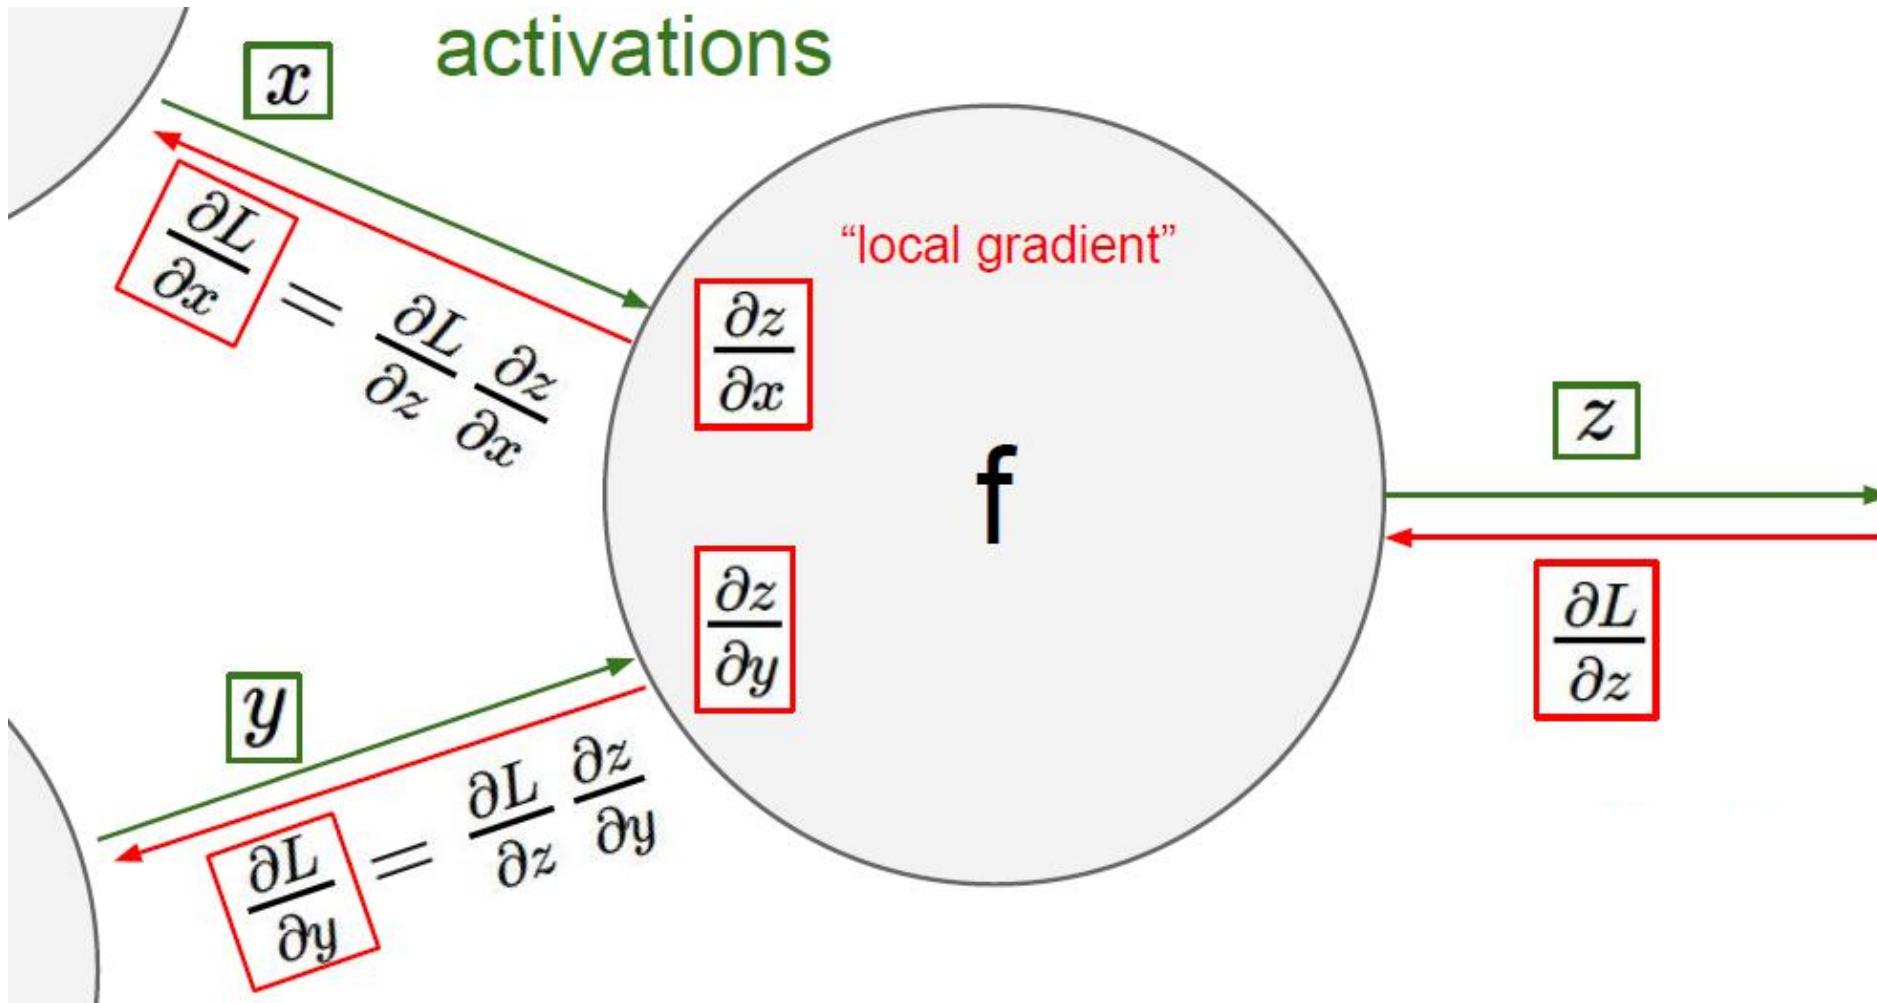
# Chain Rule



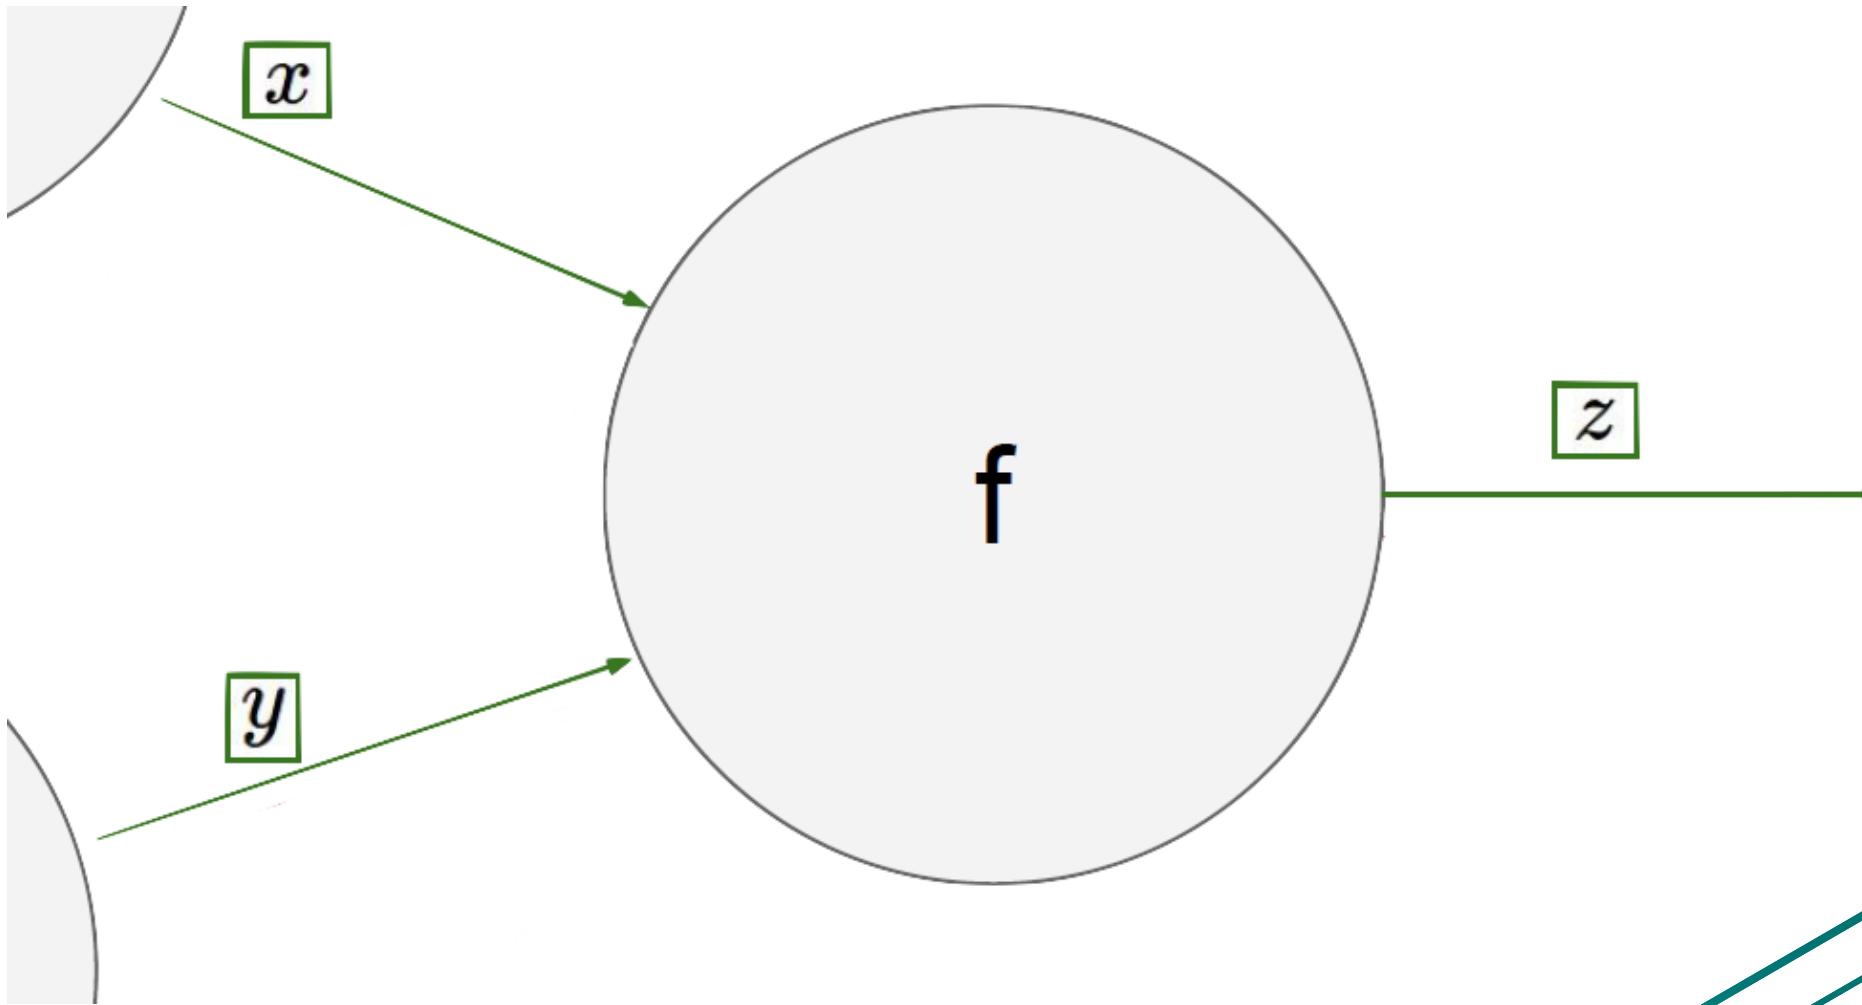
# Chain Rule



# Chain Rule

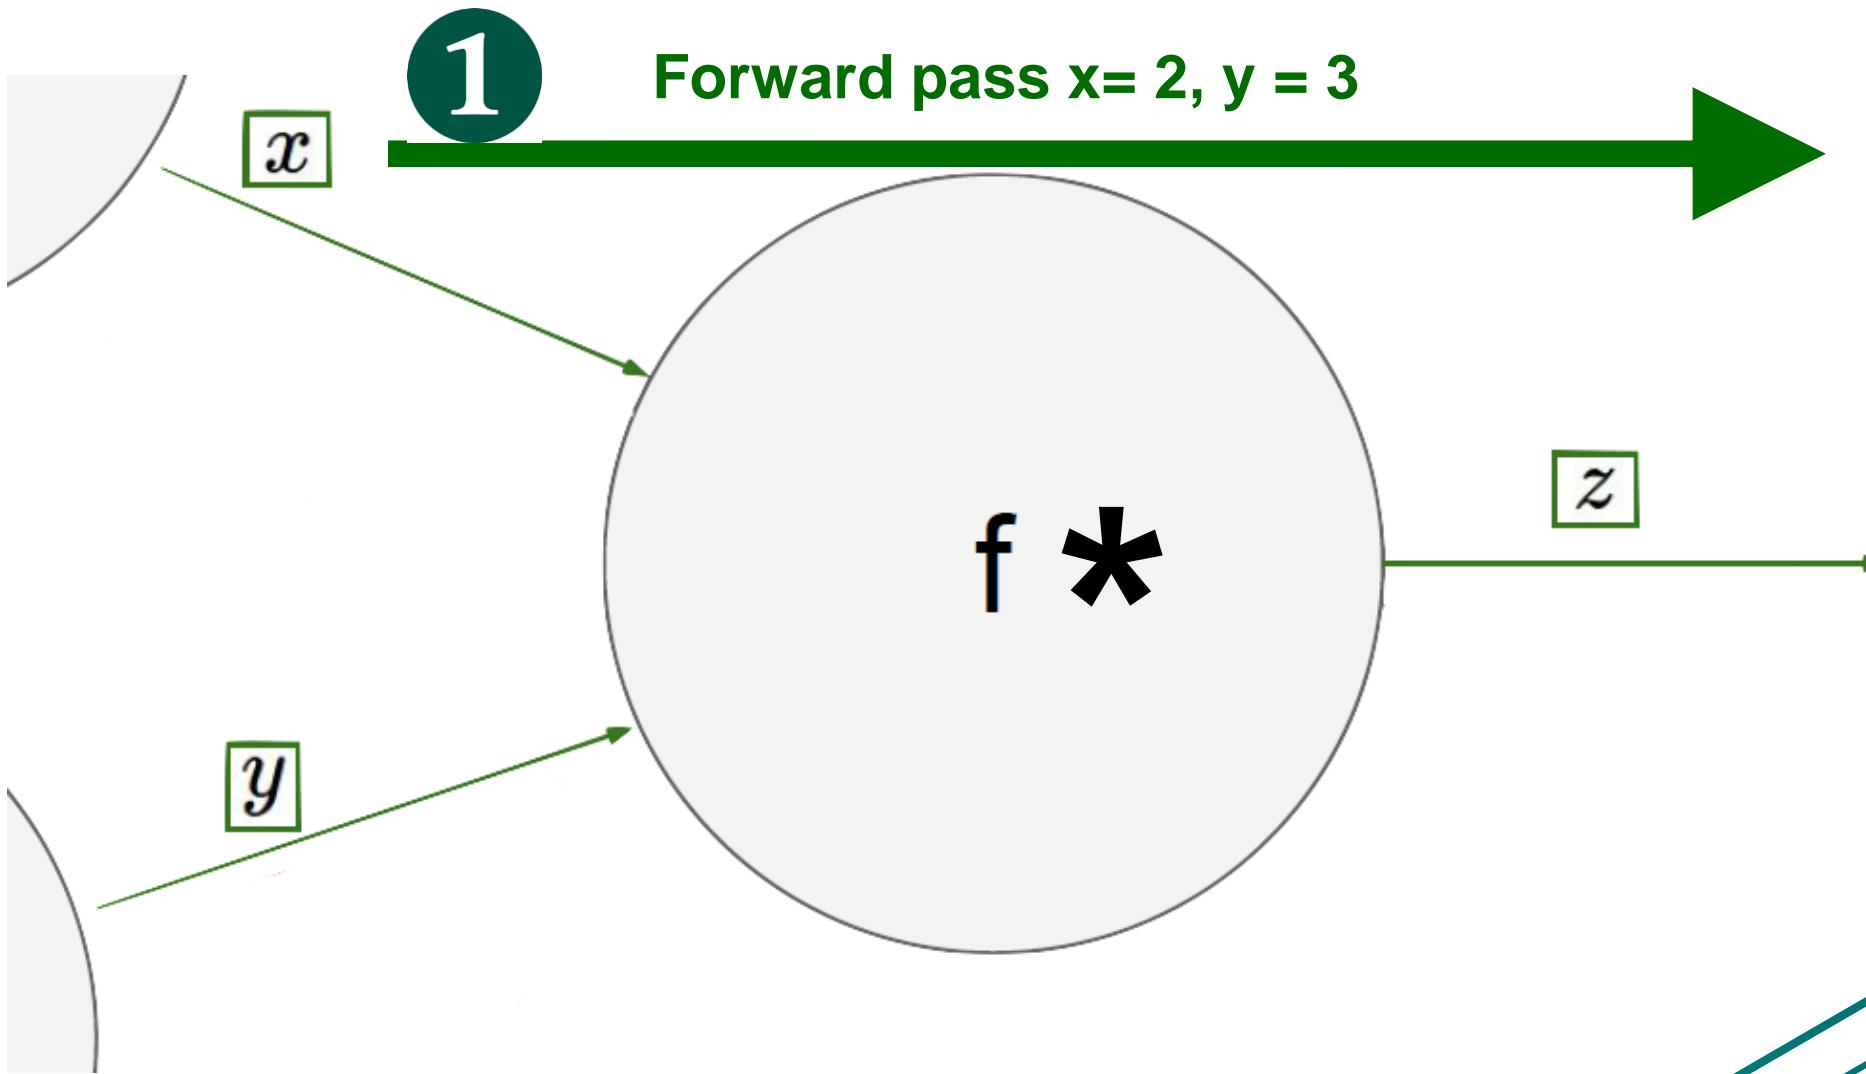


# Chain Rule



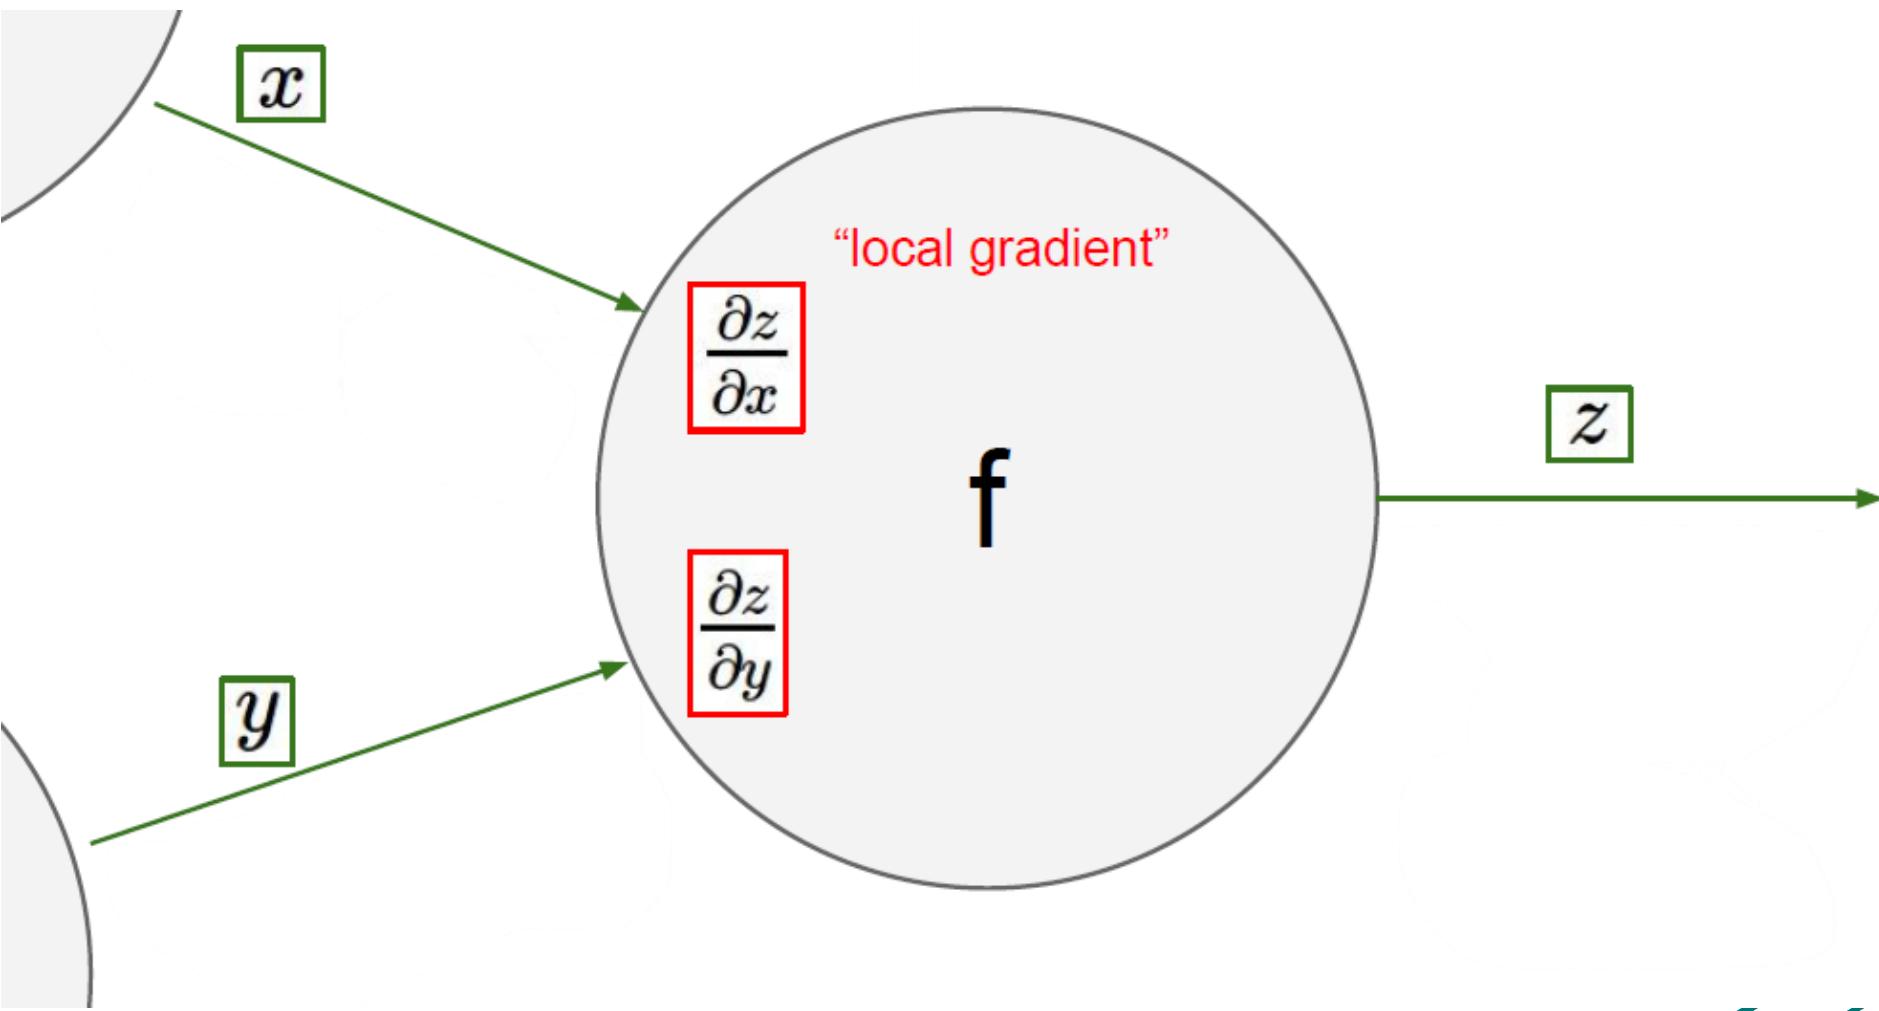
HAMIM®

# Forward pass



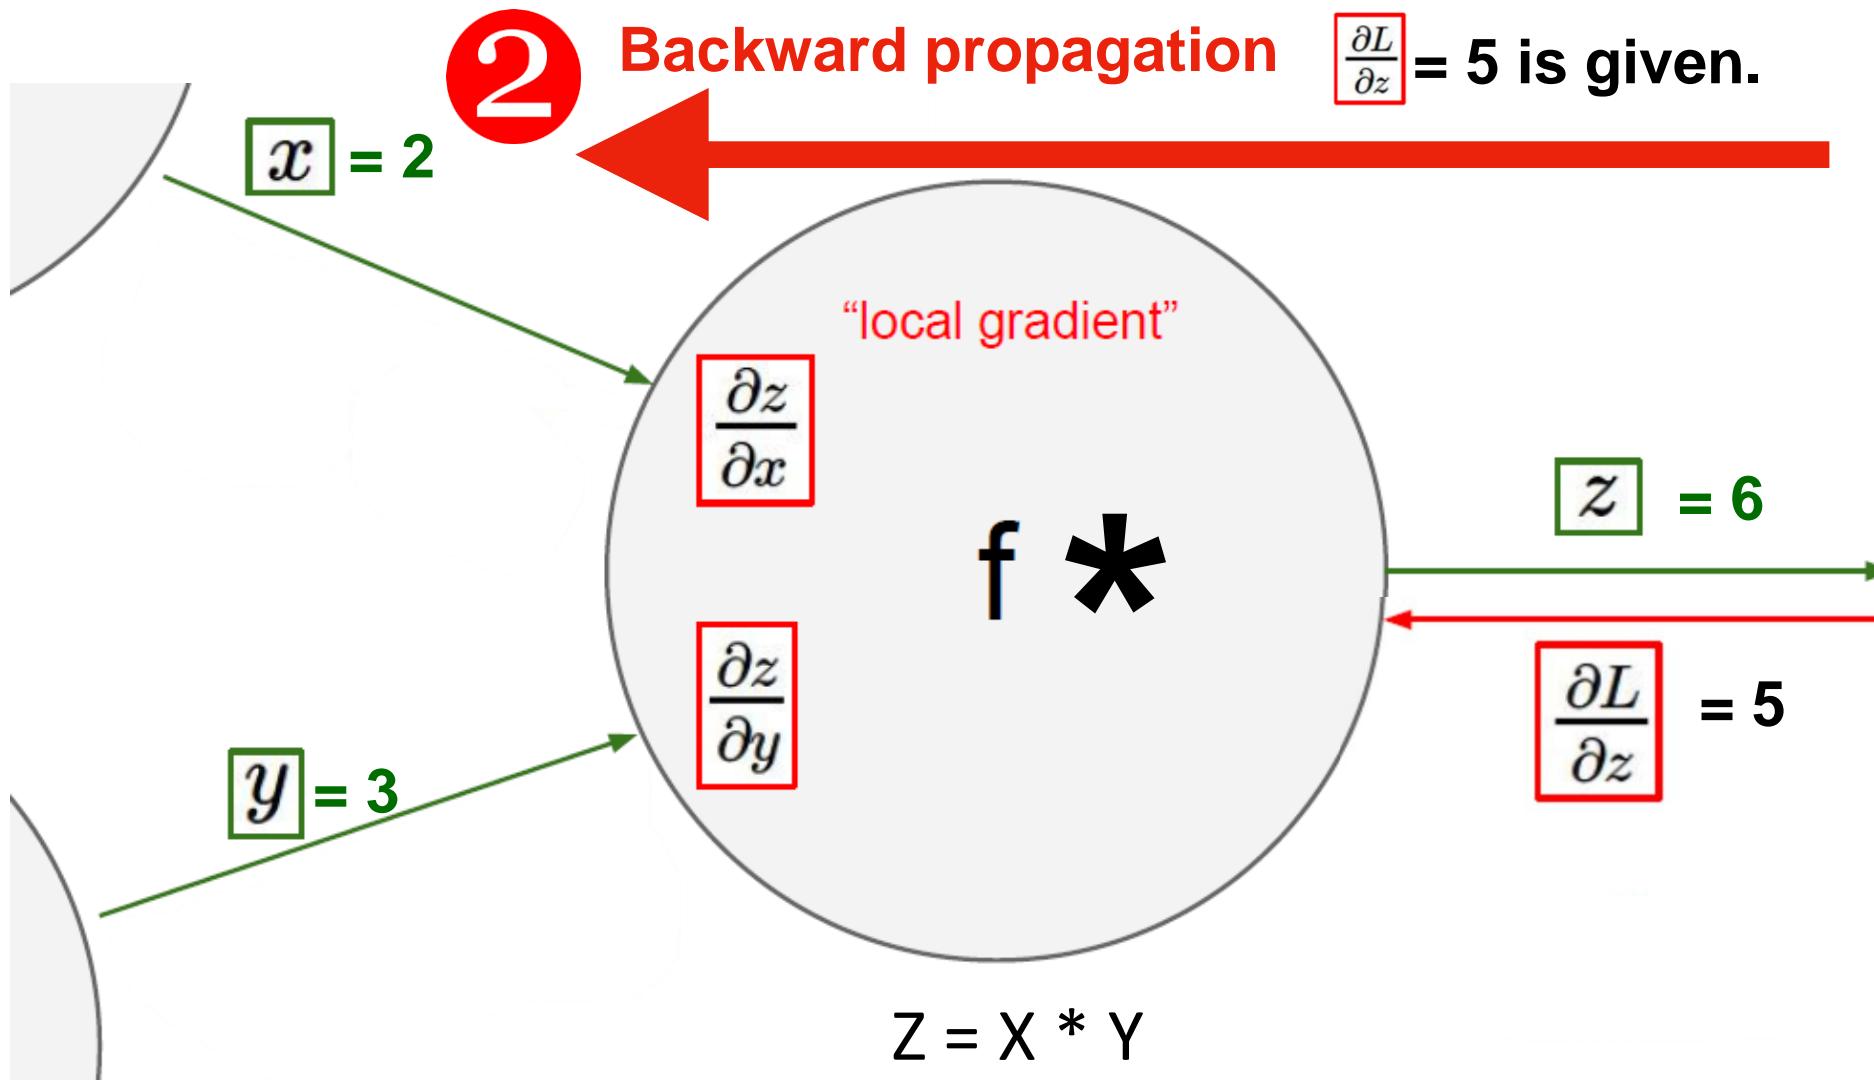
HAMIM®

# Forward pass

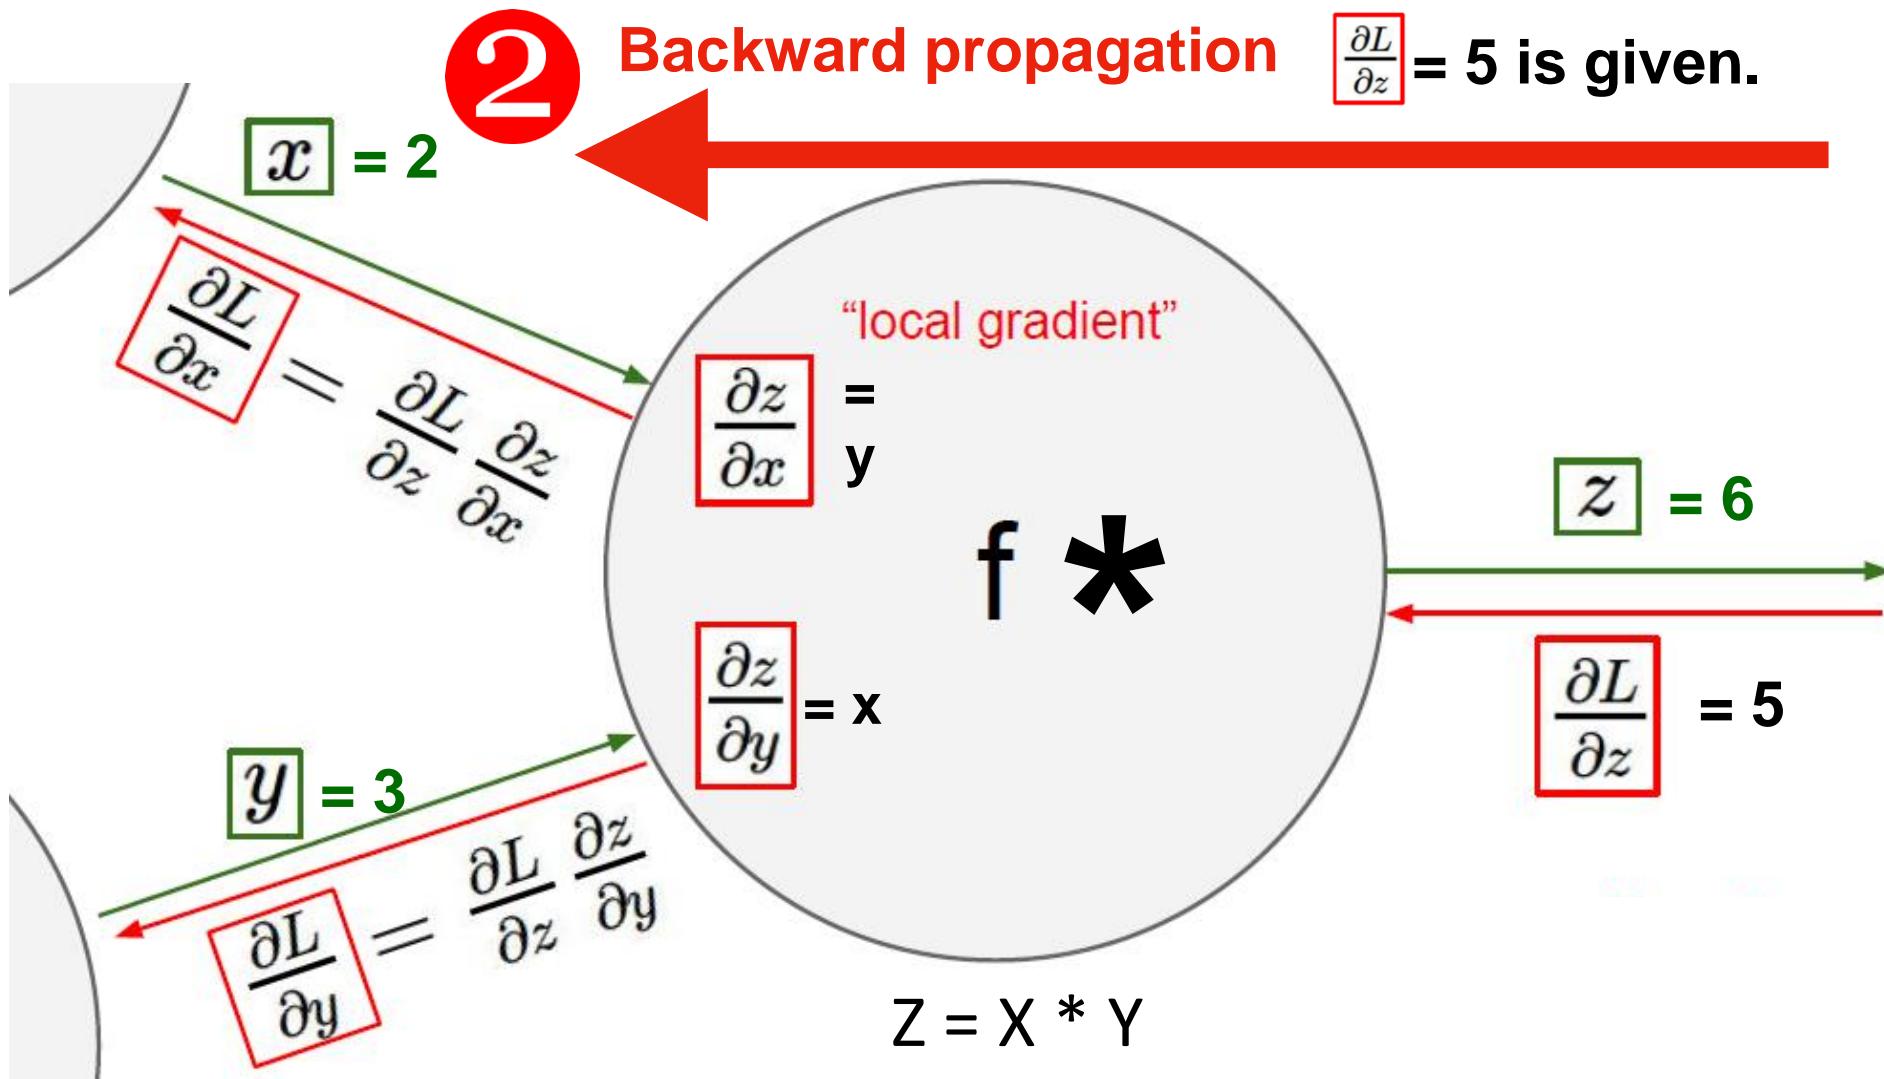


HAMIM®

# Backward propagation

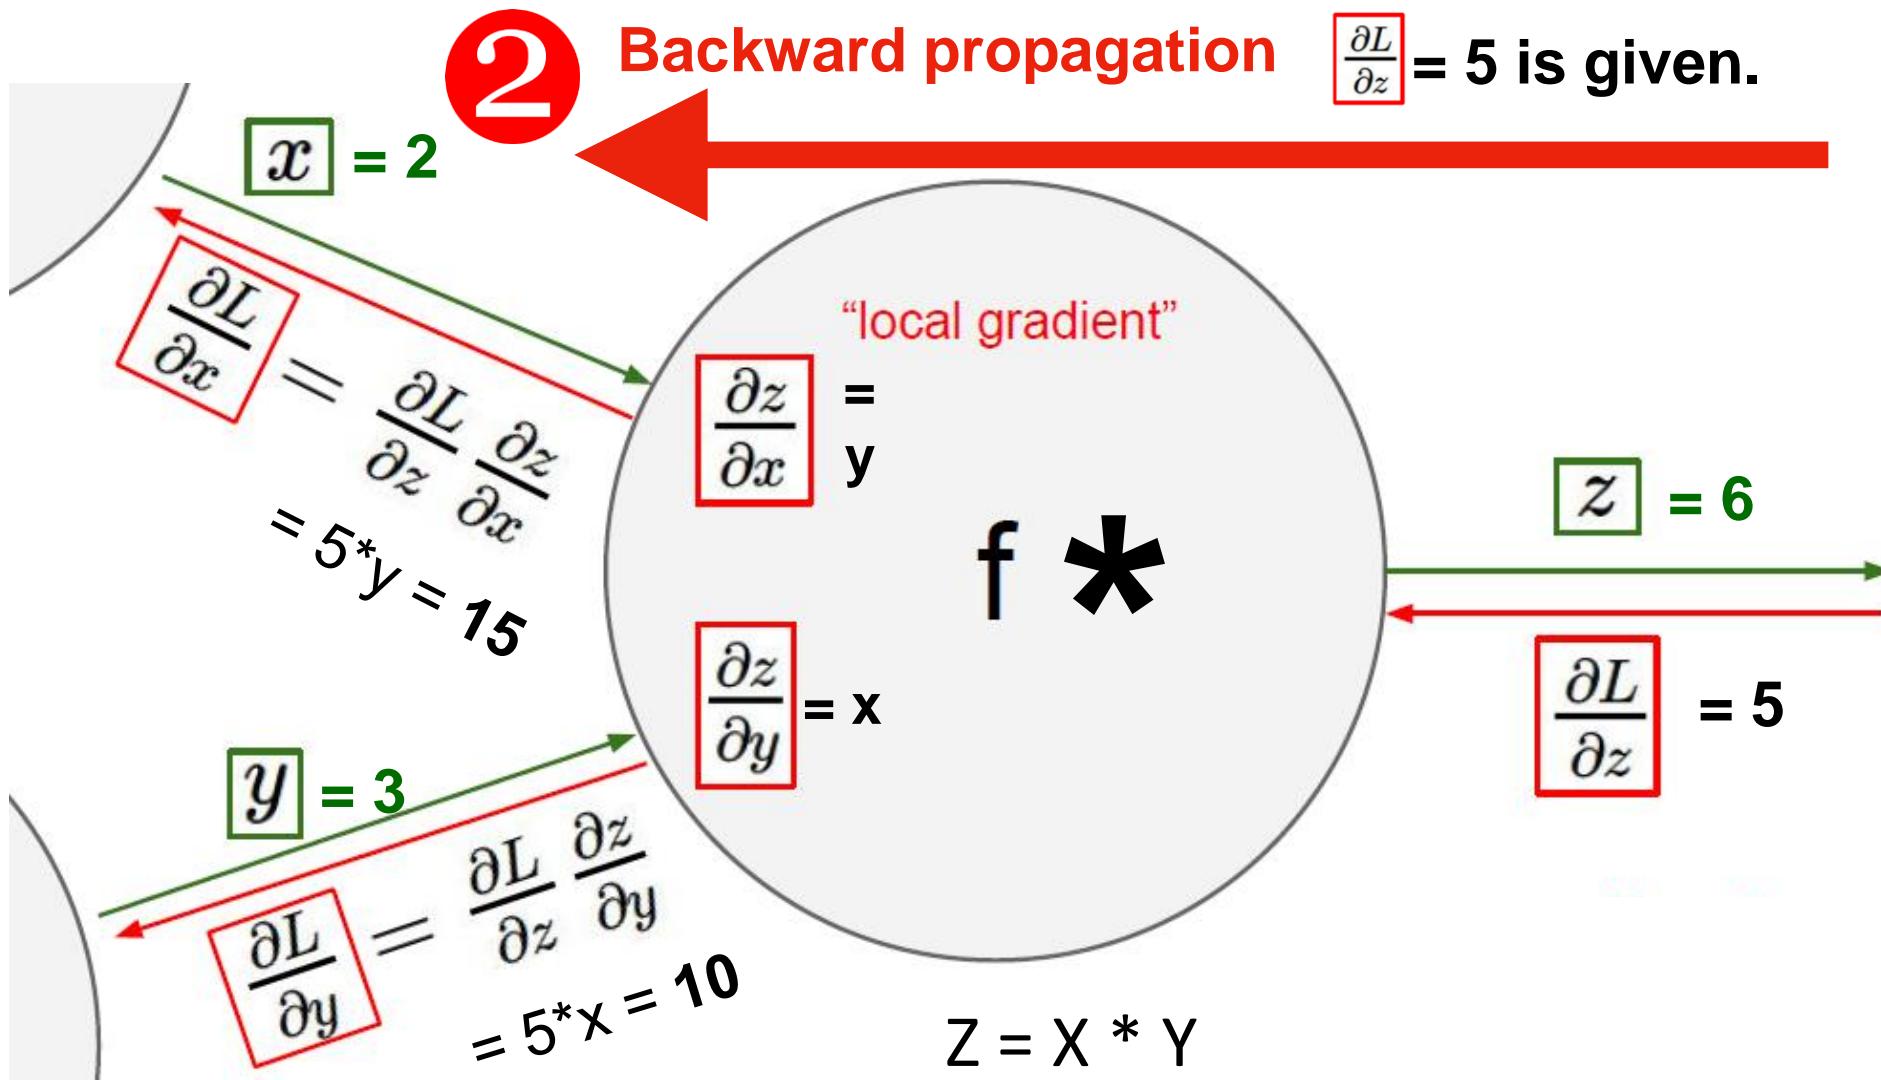


# Backward propagation



HAMIM®

# Backward Propagation



# Computational Graph

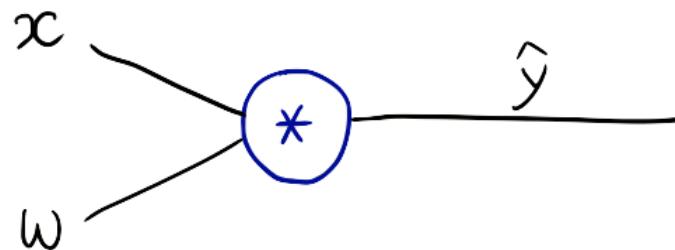
$$\hat{y} = x * w$$



HAMIM®

# Computational Graph

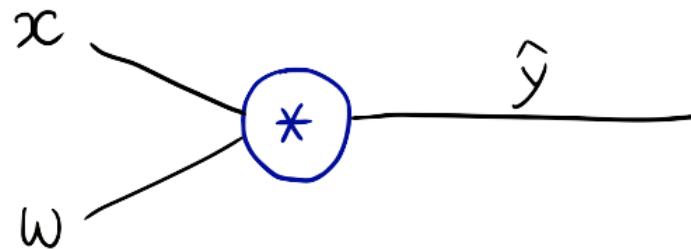
$$\hat{y} = x * w$$



HAMIM®

# Computational Graph

$$\hat{y} = x * w$$

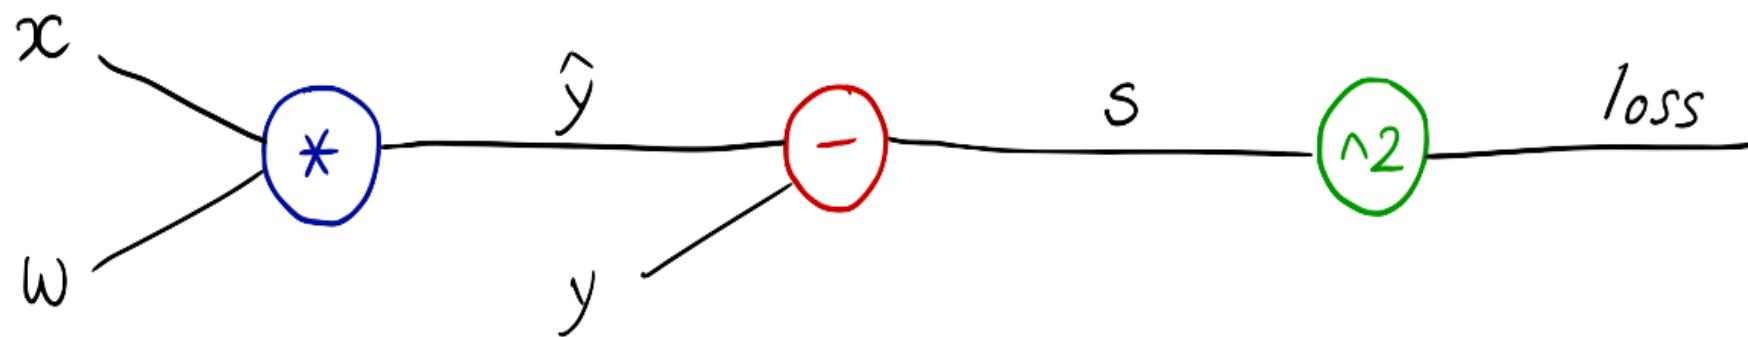


$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$



# Computational Graph

$$loss = (\hat{y} - y)^2 = (x * w - y)^2$$

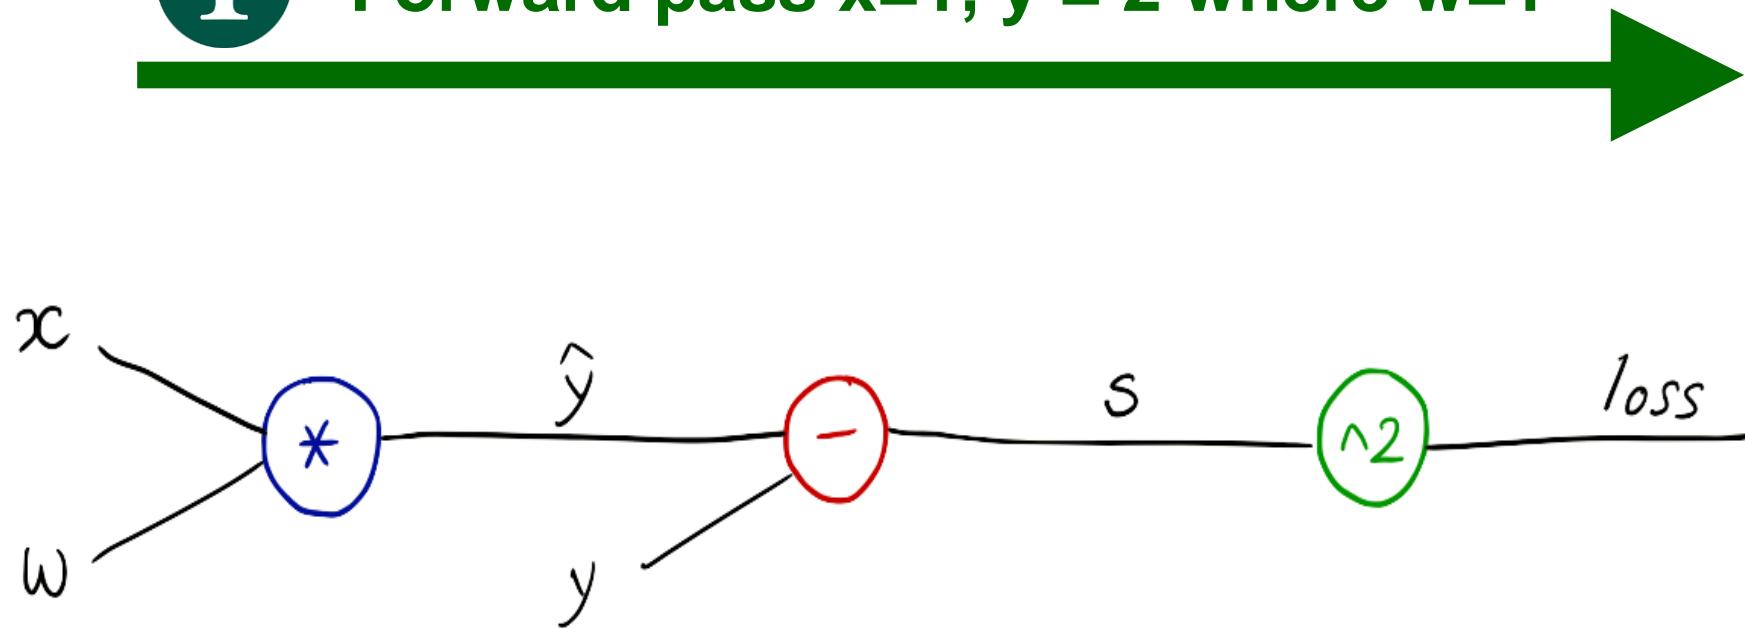


HAMIM®

# Computational Graph

1

Forward pass  $x=1, y = 2$  where  $w=1$



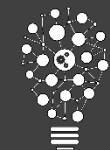
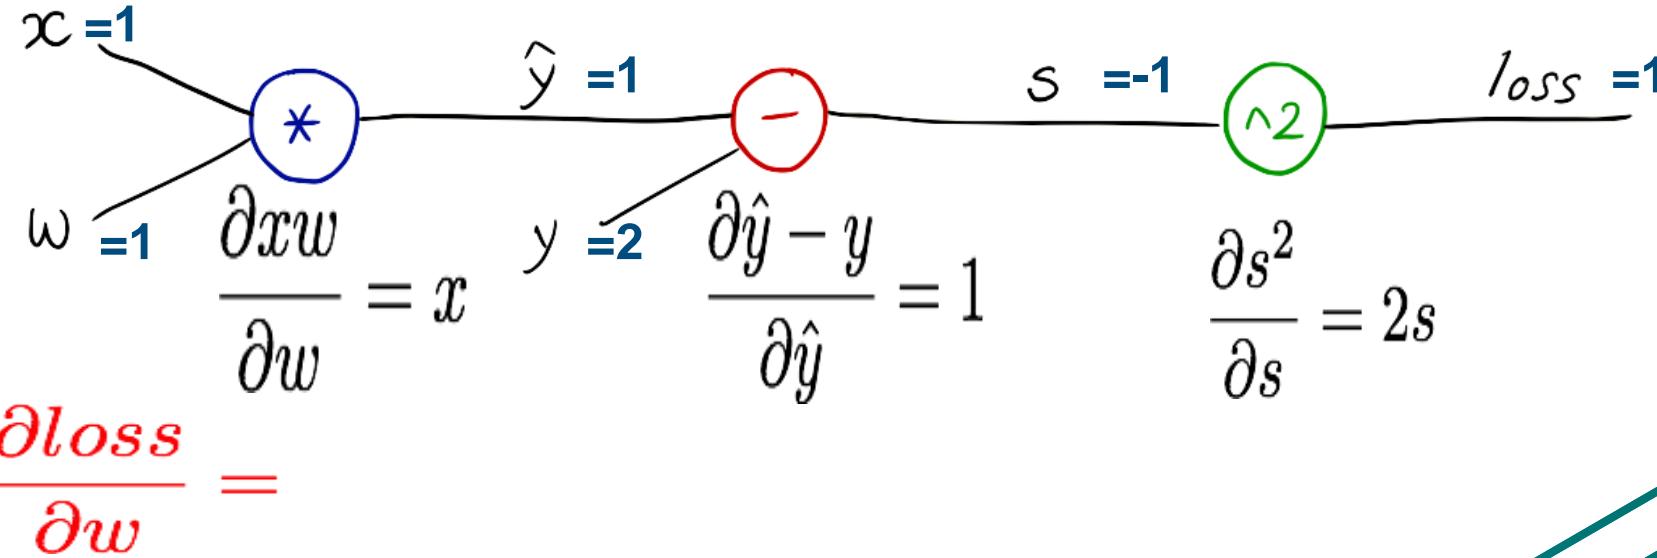
HAMIM®

# Computational Graph

$$\text{loss} = (\hat{y} - y)^2 = (x * w - y)^2$$

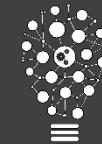
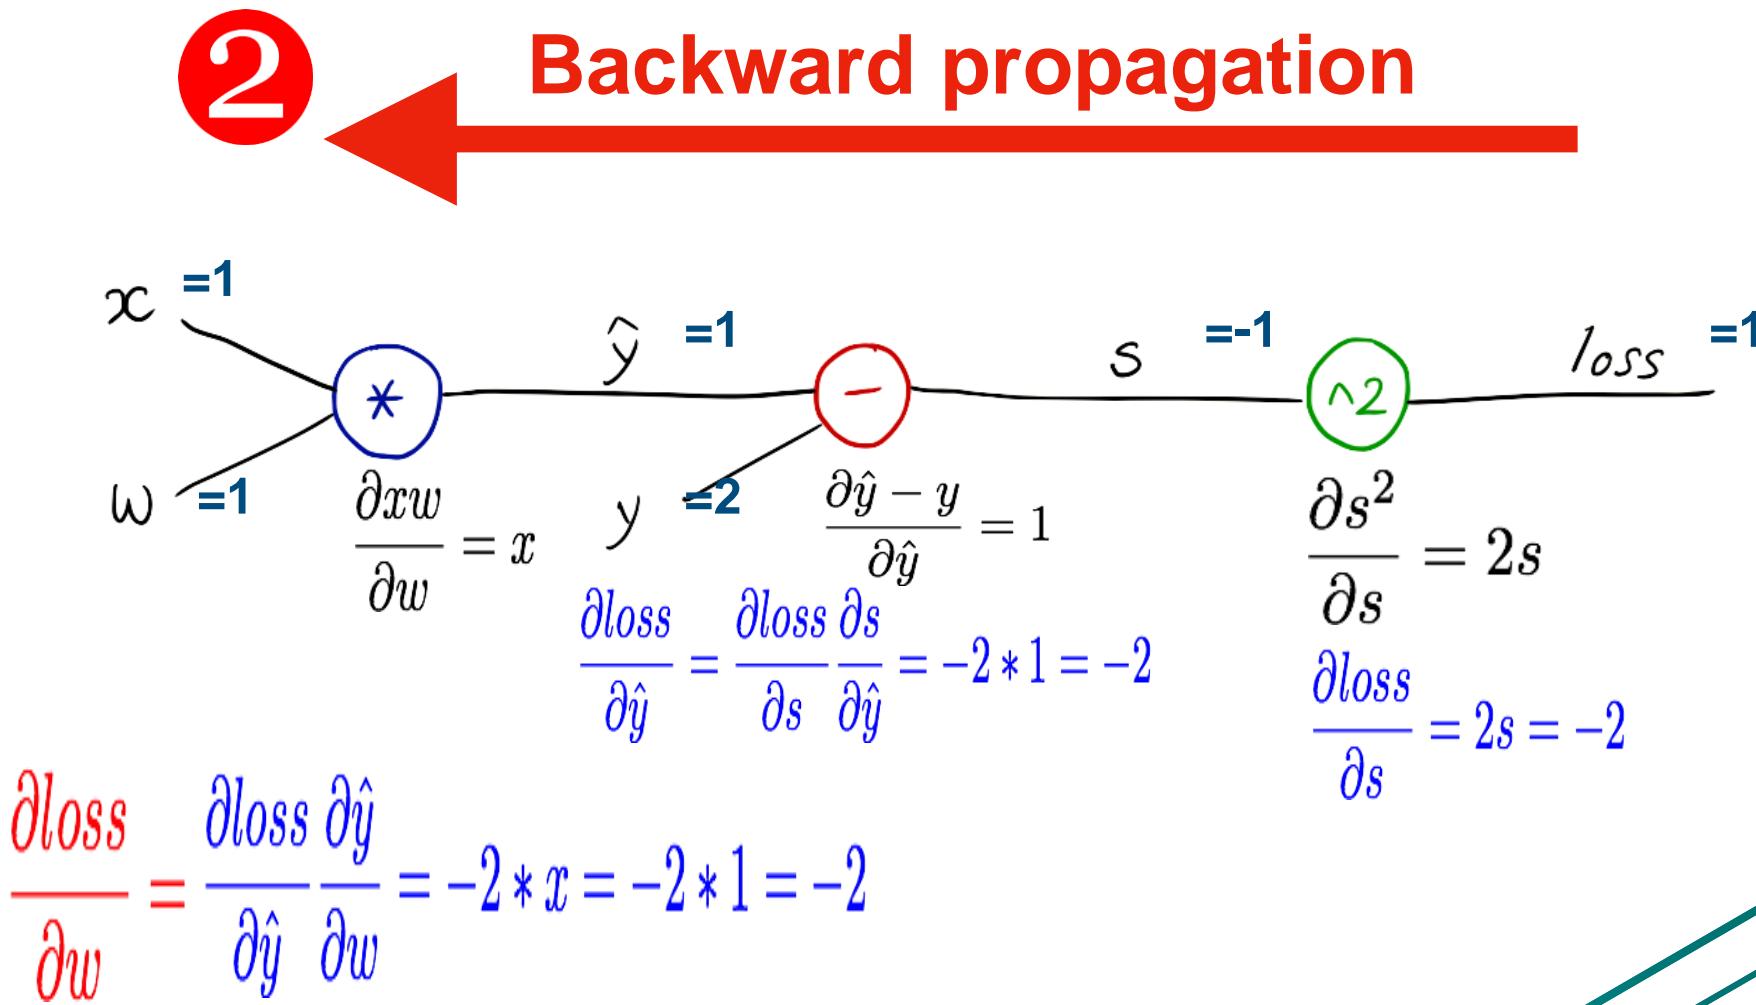
2

Backward propagation



HAMIM®

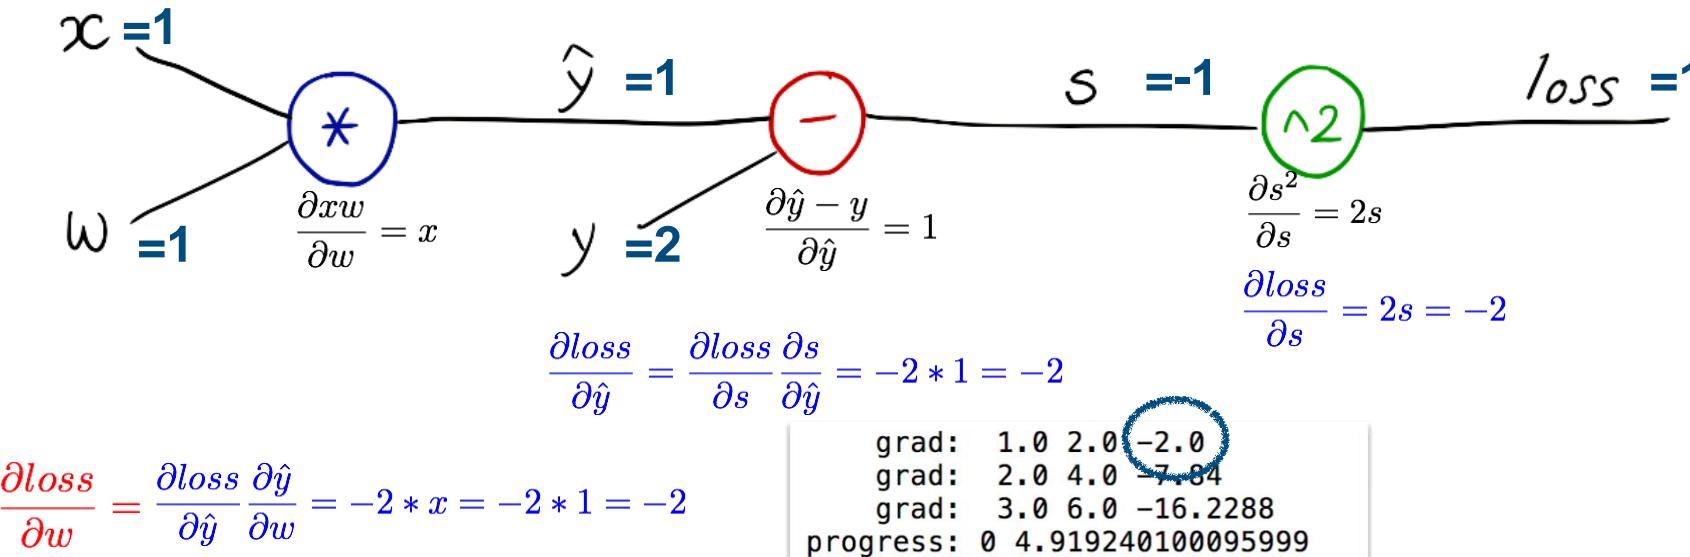
# Computational Graph



# Computational Graph

2

Backward propagation



HAMIM®



# Data and Variable

```
import torch  
from torch.autograd import Variable  
  
x_data = [1.0, 2.0, 3.0]  
y_data = [2.0, 4.0, 6.0]  
  
w = Variable(torch.Tensor([1.0]),  
            requires_grad=True) # Any random value
```

# Data and Variable



A graph is created on the fly

$W_h$        $h$        $W_x$        $x$

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```

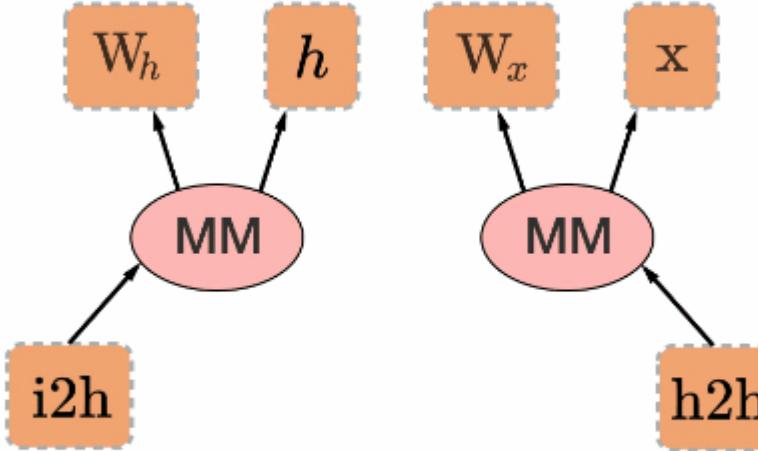
<http://pytorch.org/docs/master/notes/autograd.html?highlight=variable>

# Data and Variable



A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))  
  
i2h = torch.mm(W_x, x.t())  
h2h = torch.mm(W_h, prev_h.t())
```



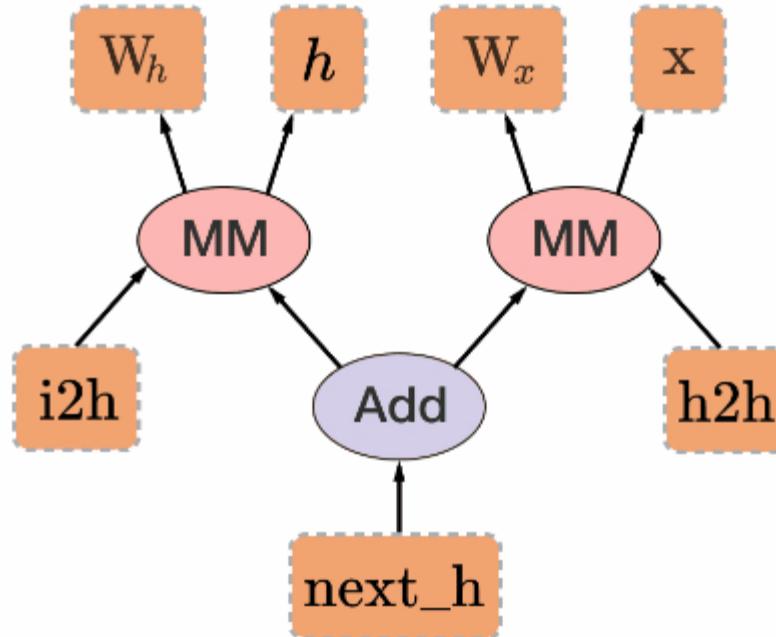
<http://pytorch.org/docs/master/notes/autograd.html?highlight=variable>

# Data and Variable



A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))  
  
i2h = torch.mm(W_x, x.t())  
h2h = torch.mm(W_h, prev_h.t())  
next_h = i2h + h2h
```



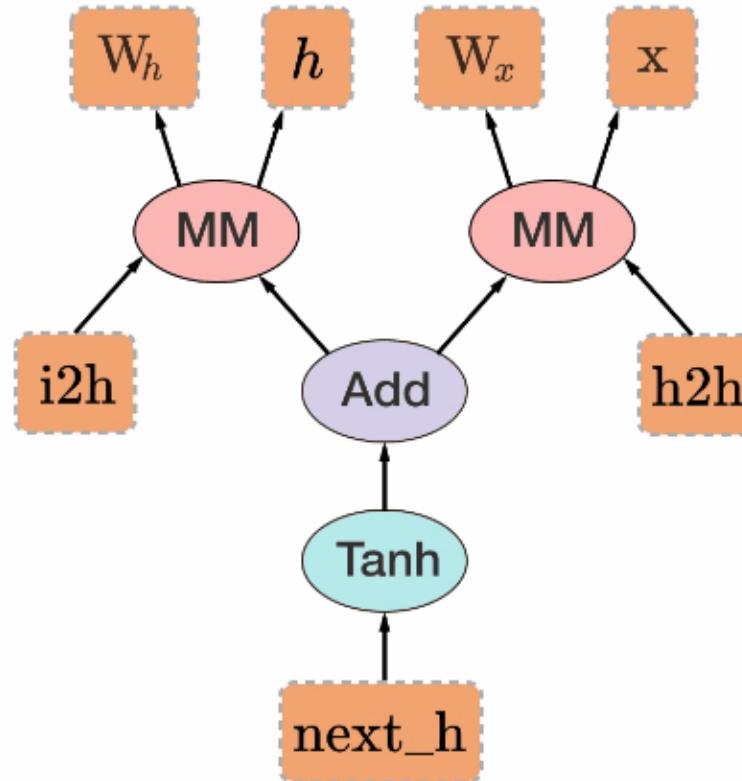
<http://pytorch.org/docs/master/notes/autograd.html?highlight=variable>

# Data and Variable



A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))  
  
i2h = torch.mm(W_x, x.t())  
h2h = torch.mm(W_h, prev_h.t())  
next_h = i2h + h2h  
next_h = next_h.tanh()
```



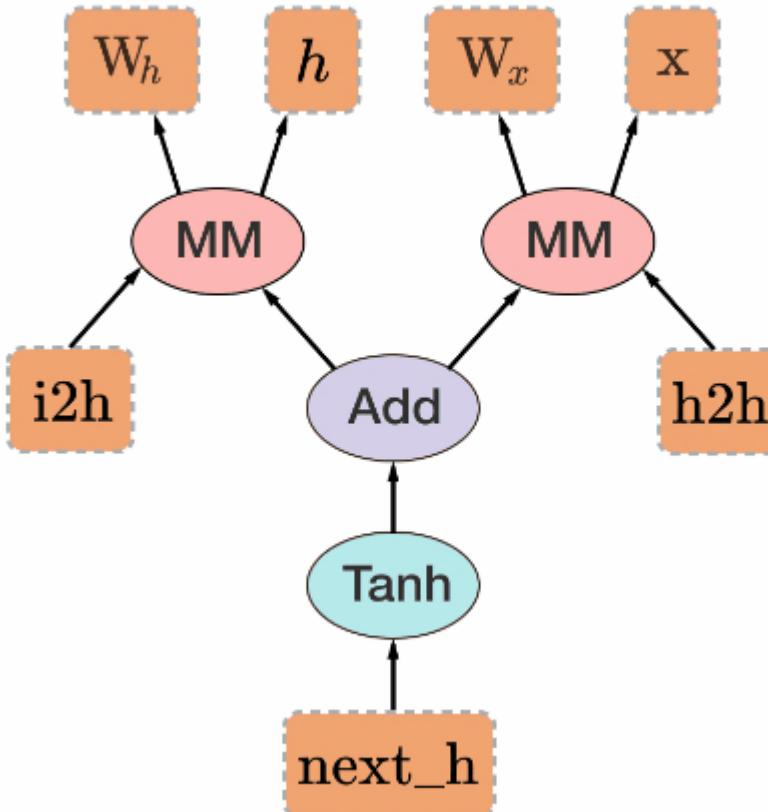
<http://pytorch.org/docs/master/notes/autograd.html?highlight=variable>

# Data and Variable



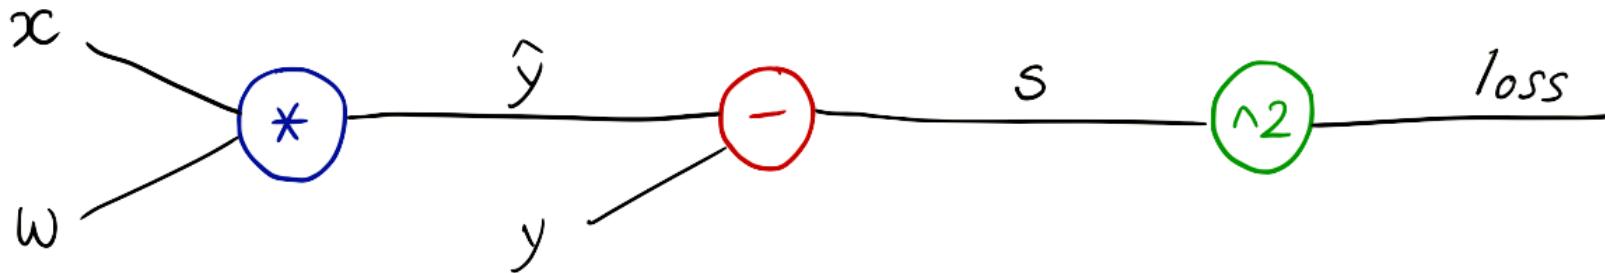
Back-propagation  
uses the dynamically built graph

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))  
  
i2h = torch.mm(W_x, x.t())  
h2h = torch.mm(W_h, prev_h.t())  
next_h = i2h + h2h  
next_h = next_h.tanh()  
  
next_h.backward(torch.ones(1, 20))
```



<http://pytorch.org/docs/master/notes/autograd.html?highlight=variable>

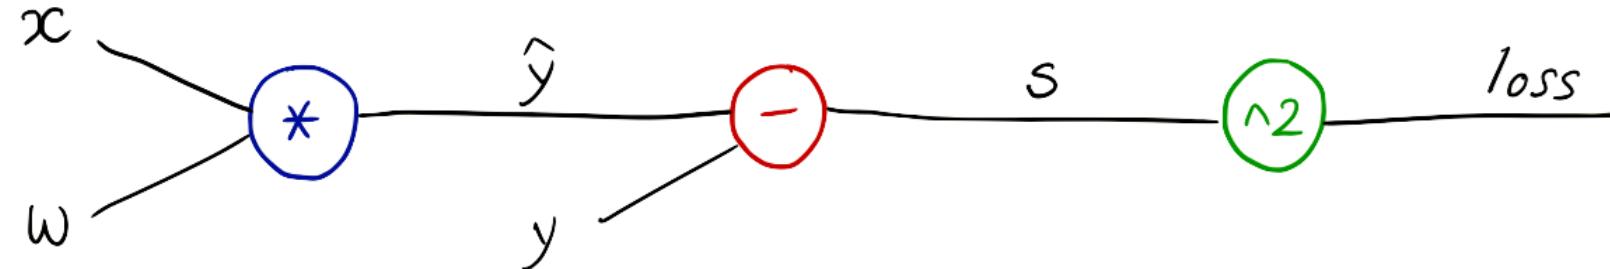
# PyTorch forward/backward



HAMIM®

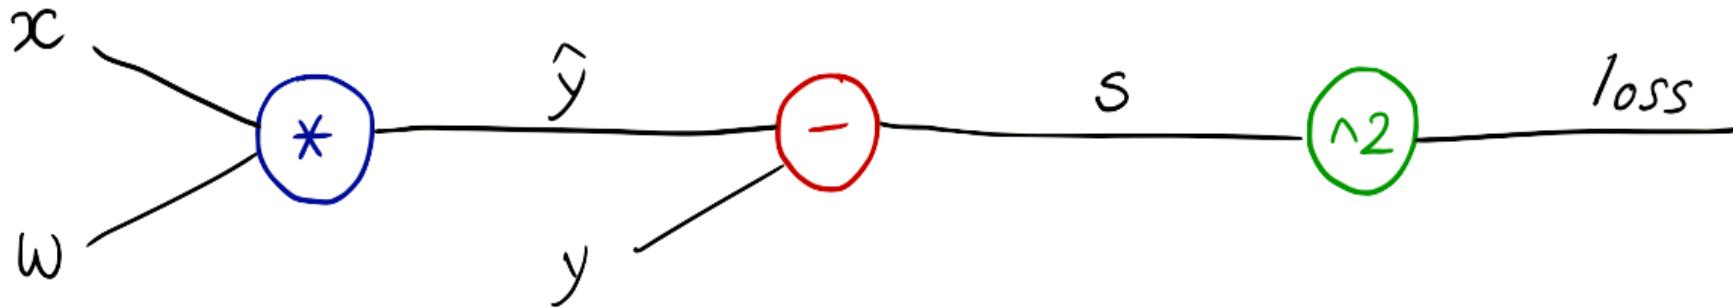
# Forward Pass

```
# Any random value  
w = Variable(torch.Tensor([1.0]), requires_grad=True)  
l = loss(x=1, y= )
```



$$\frac{\partial \text{loss}}{\partial w} =$$

# Back propagation : l.backward()

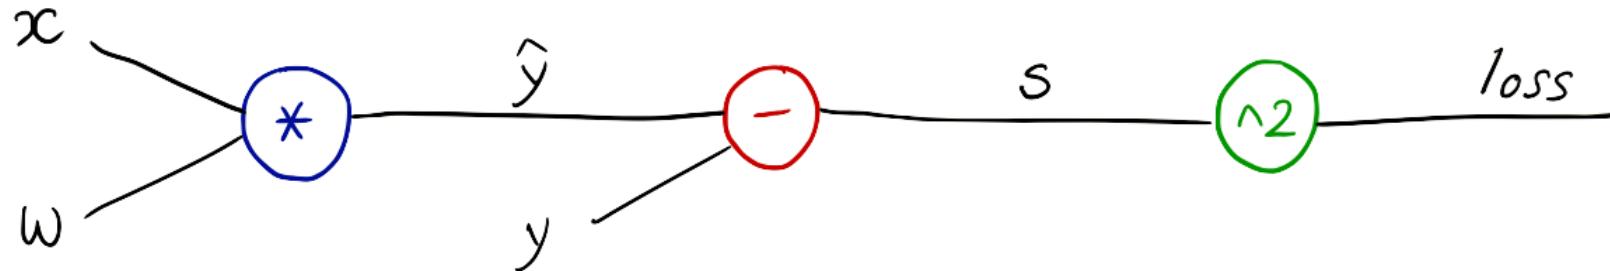


$$\frac{\partial \text{loss}}{\partial w} = \text{w.grad}$$

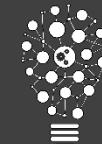


# Weight update (step)

```
w.data = w.data - 0.01 * w.grad.data
```



$$\frac{\partial loss}{\partial w} = w.grad$$





# Model and Loss

```
# our model forward pass
def forward(x):
    return x * w

# Loss function
def loss(x, y):
    y_pred = forward(x)
    return (y_pred - y) * (y_pred - y)

# Before training
print("predict (before training)", 4, forward(4).data[0])
```



HAMIM®

# Training: forward, backward, Update Weight



```
# Training Loop
for epoch in range(10):
    for x_val, y_val in zip(x_data, y_data):
        l = loss(x_val, y_val)
        l.backward()
        print("\tgrad: ", x_val, y_val, w.grad.data[0])
        w.data = w.data - 0.01 * w.grad.data

    # Manually zero the gradients after updating weights
    w.grad.data.zero_()

    print("progress:", epoch, l.data[0])

# After training
print("predict (after training)", 4, forward(4).data[0])
```

# Output

```
# Training Loop
for epoch in range(10):
    for x_val, y_val in zip(x_data, y_data):
        l = loss(x_val, y_val)
        l.backward()
        print("\tgrad: ", x_val, y_val, w.grad.data[0])
        w.data = w.data - 0.01 * w.grad.data

    # Manually zero the gradients after updating weights
    w.grad.data.zero_()

    print("progress:", epoch, l.data[0])

# After training
print("predict (after training)", 4, forward(4).data[0])
```



```
predict (before training) 4 4.0
grad: 1.0 2.0 -2.0
grad: 2.0 4.0 -7.840000152587891
grad: 3.0 6.0 -16.228801727294922
progress: 0 7.315943717956543
grad: 1.0 2.0 -1.478623867034912
grad: 2.0 4.0 -5.796205520629883
grad: 3.0 6.0 -11.998146057128906
progress: 1 3.9987640380859375
grad: 1.0 2.0 -1.0931644439697266
grad: 2.0 4.0 -4.285204887390137
grad: 3.0 6.0 -8.870372772216797
progress: 2 2.1856532096862793
grad: 1.0 2.0 -0.8081896305084229
grad: 2.0 4.0 -3.1681032180786133
grad: 3.0 6.0 -6.557973861694336
progress: 3 1.1946394443511963
grad: 1.0 2.0 -0.5975041389465332
grad: 2.0 4.0 -2.3422164916992188
grad: 3.0 6.0 -4.848389625549316
progress: 4 0.6529689431190491
grad: 1.0 2.0 -0.4417421817779541
grad: 2.0 4.0 -1.7316293716430664
grad: 3.0 6.0 -3.58447265625
progress: 5 0.35690122842788696
grad: 1.0 2.0 -0.3265852928161621
grad: 2.0 4.0 -1.2802143096923828
grad: 3.0 6.0 -2.650045394897461
progress: 6 0.195076122879982
grad: 1.0 2.0 -0.24144840240478516
grad: 2.0 4.0 -0.9464778900146484
grad: 3.0 6.0 -1.9592113494873047
progress: 7 0.10662525147199631
grad: 1.0 2.0 -0.17850565910339355
grad: 2.0 4.0 -0.699742317199707
grad: 3.0 6.0 -1.4484672546386719
```



# Output (from numeric & gradient Descent)

```
predict (before training) 4 4.0
grad: 1.0 2.0 -2.0
grad: 2.0 4.0 -7.84
grad: 3.0 6.0 -16.2288
progress: 0 4.919240100095999
grad: 1.0 2.0 -1.478624
grad: 2.0 4.0 -5.796206079999999
grad: 3.0 6.0 -11.998146585599997
progress: 1 2.688769240265834
grad: 1.0 2.0 -1.093164466688
grad: 2.0 4.0 -4.285204709416961
grad: 3.0 6.0 -8.87037374849311
progress: 2 1.4696334962911515
grad: 1.0 2.0 -0.8081896081960389
grad: 2.0 4.0 -3.1681032641284723
grad: 3.0 6.0 -6.557973756745939
progress: 3 0.8032755585999681
grad: 1.0 2.0 -0.59750427561463
grad: 2.0 4.0 -2.3422167604093502
grad: 3.0 6.0 -4.848388694047353
progress: 4 0.43905614881022015
grad: 1.0 2.0 -0.44174208101320334
grad: 2.0 4.0 -1.7316289575717576
grad: 3.0 6.0 -3.584471942173538
progress: 5 0.2399802903801062
grad: 1.0 2.0 -0.3265852213980338
grad: 2.0 4.0 -1.2802140678802925
grad: 3.0 6.0 -2.650043120512205
progress: 6 0.1311689630744999
grad: 1.0 2.0 -0.241448373202223
grad: 2.0 4.0 -0.946477622952715
grad: 3.0 6.0 -1.9592086795121197
progress: 7 0.07169462478267678
grad: 1.0 2.0 -0.17850567968888198
grad: 2.0 4.0 -0.6997422643804168
grad: 3.0 6.0 -1.4484664872674653
progress: 8 0.03918700813247573
grad: 1.0 2.0 -0.13197139106214673
grad: 2.0 4.0 -0.5173278529636143
grad: 3.0 6.0 -1.0708686556346834
progress: 9 0.021418922423117836
predict (after training) 4 7.804863933862125
```

```
# Before training
print("predict (before training)", 4, forward(4))

# Training loop
for epoch in range(10):
    for x, y in zip(x_data, y_data):
        grad = gradient(x, y)
        w = w - 0.01 * grad
        print("\tgrad: ", x, y, grad)
        l = loss(x, y)

    print ("progress:", epoch, l)

# After training
print("predict (after training)", 4, forward(4))
```



# Classification Problem

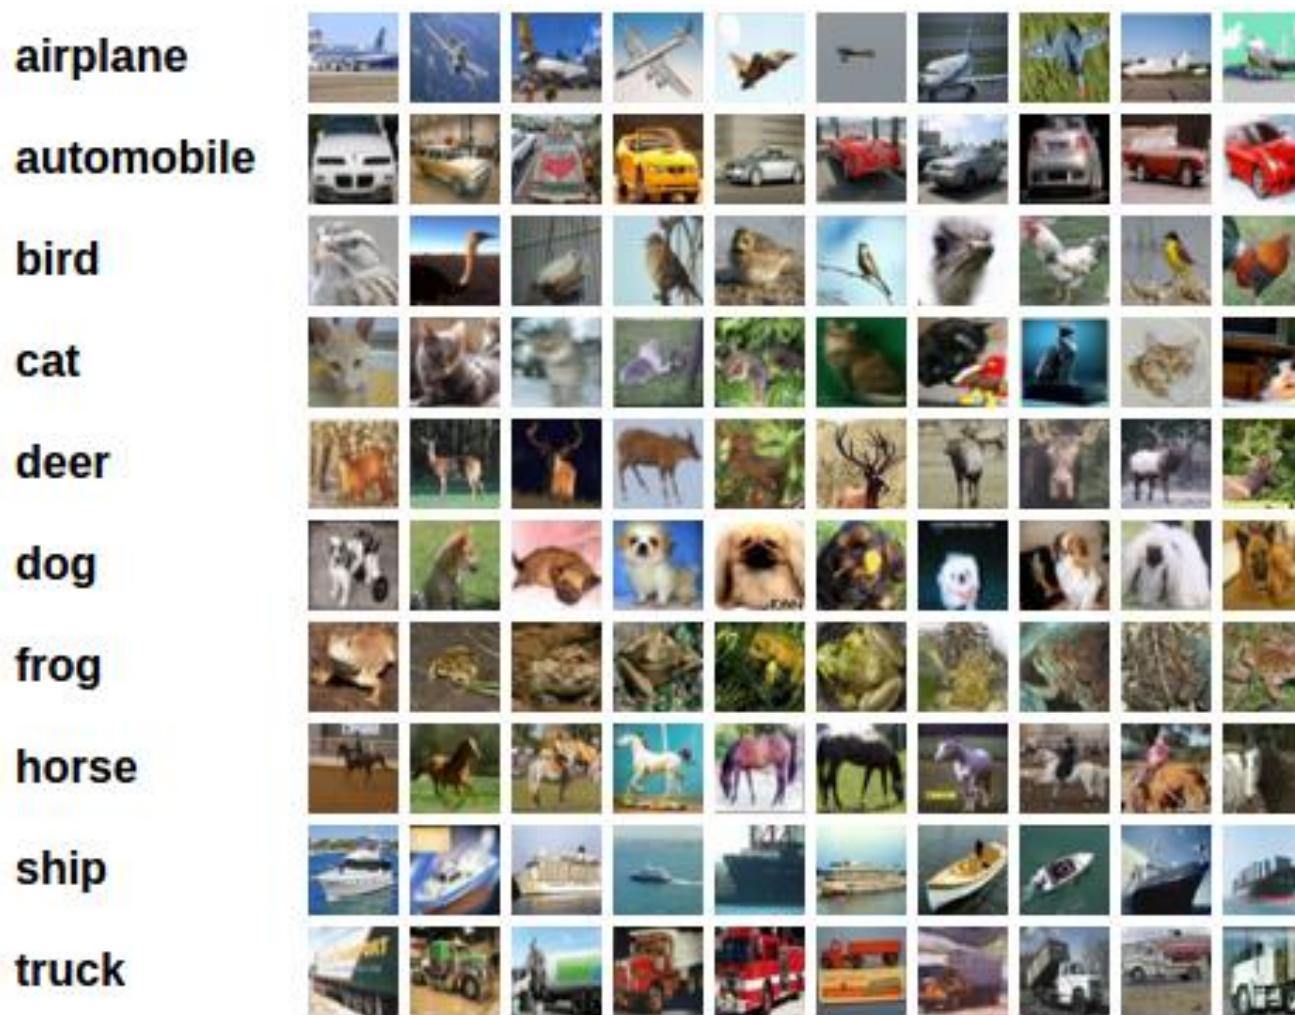


Image Size =  $3 \times 32 \times 32$

# PyTorch Rhythm

PYTORCH

- 1 Design your model using class with Variables
- 2 Construct loss and optimizer (select from PyTorch API)
- 3 Training cycle (forward, backward, update)



HAMIM®

# Classification Problem

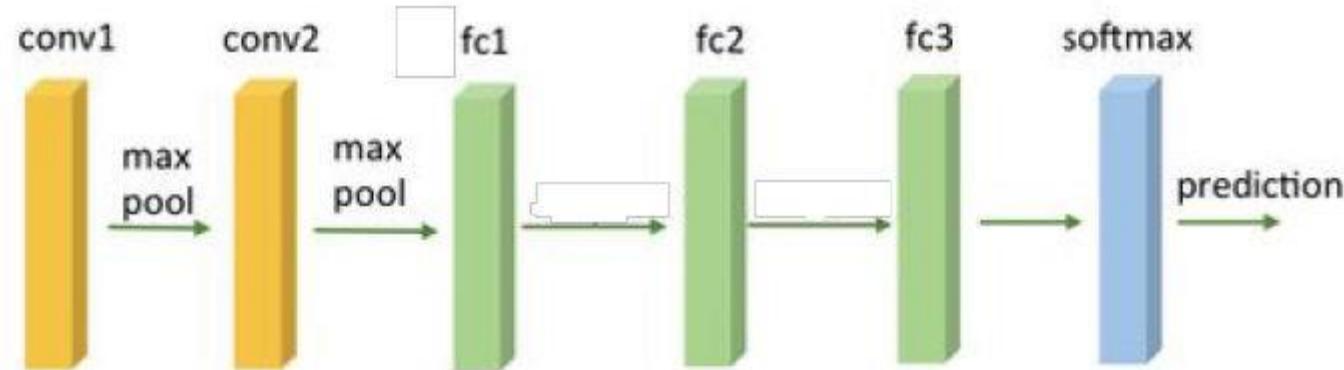
## 1 Design your model using class with Variables

```
#####
# 1. Define a Neural Network
# ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
# Copy the neural network from the Neural Networks section before and modify it to
# take 3-channel images (instead of 1-channel images as it was defined).
```

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
net = Net()
```



# Cifar10 classification

## ② Construct loss and optimizer (select from PyTorch API)

```
#####
# 2. Define a Loss function and optimizer
# ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
# Let's use a Classification Cross-Entropy loss and SGD with momentum
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

#####
```

# Cifar10 classification

```
#####
# 3. Train the network
# ~~~~~
#
# This is when things start to get interesting.
# We simply have to loop over our data iterator, and feed the inputs to the
# network and optimize
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data

        # wrap them in Variable
        inputs, labels = Variable(inputs), Variable(labels)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.data[0]
        if i % 2000 == 1999: # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

print('Finished Training')
```

## 3 Training cycle (forward, backward, update)



HAMIM®

# Torch.nn for Convolution layers structure

- `m = nn.Conv2d(16, 33, 3, stride=2, padding=3)`
- Parameters:
  - `in_channels` (int) – Number of channels in the input image
  - `out_channels` (int) – Number of channels produced by the convolution
  - `kernel_size` (int or tuple) – Size of the convolving kernel
  - `stride` (int or tuple, optional) – Stride of the convolution. Default: 1
  - `padding` (int or tuple, optional) – Zero-padding added to both sides of the input. Default: 0

## Other Layers

- Convolution Layers
- Pooling Layers
- Padding Layers
- Non-linear Activations
- Normalization Layers
- Recurrent layers
- Linear Layers
- Dropout Layers
- Sparse layers
- Vision layers

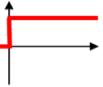
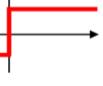
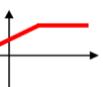
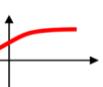
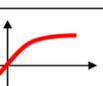
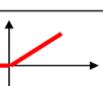
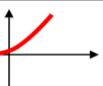


HAMIM®

# Torch Optimization Function

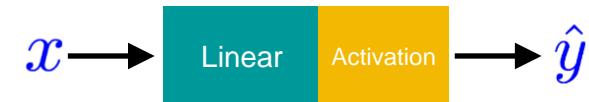
- `torch.optim.Adagrad`
- `torch.optim.Adam`
- `torch.optim.Adamax`
- `torch.optim.ASGD`
- `torch.optim.LBFGS`
- `torch.optim.RMSprop`
- `torch.optim.Rprop`
- `torch.optim.SGD`

# Activation Functions

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer Neural Networks	
Rectifier, ReLU (Rectified Linear Unit)	$\phi(z) = \max(0, z)$	Multi-layer Neural Networks	
Rectifier, softplus	$\phi(z) = \ln(1 + e^z)$	Multi-layer Neural Networks	

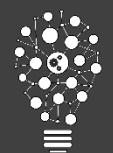
Copyright © Sebastian Raschka 2016  
(<http://sebastianraschka.com>)

[http://rasbt.github.io/mlxtend/user\\_guide/general\\_concepts/activation-functions/](http://rasbt.github.io/mlxtend/user_guide/general_concepts/activation-functions/)



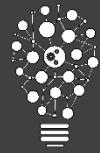
## Non-linear Activations

- ReLU
- ReLU6
- ELU
- SELU
- PReLU
- LeakyReLU
- Threshold
- Hardtanh
- Sigmoid
- Tanh
- LogSigmoid
- Softplus
- Softshrink
- Softsign
- Tanhshrink
- Softmin
- Softmax
- Softmax2d
- LogSoftmax



HAMIM®

jupyter



HAMIM®

# Data Loader \_ Manual Data Feed

```
xy = np.loadtxt('data-diabetes.csv', delimiter=',',
dtype=np.float32)
x_data = Variable(torch.from_numpy(xy[:, 0:-1]))
y_data = Variable(torch.from_numpy(xy[:, [-1]]))

# Training Loop
for epoch in range(100):
    # Forward pass: Compute predicted y by passing x to the
    model
    y_pred = model(x_data)
    # Compute and print loss
    loss = criterion(y_pred, y_data)
    print(epoch, loss.data[0])
    # Zero gradients, perform a backward pass, and update the
    weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```



# Batch (batch size)

```
# Training cycle  
for epoch in range(training_epochs):  
    # Loop over all batches  
    for i in range(total_batch):  
        batch_xs, batch_ys = ...
```



In the neural network terminology:

288

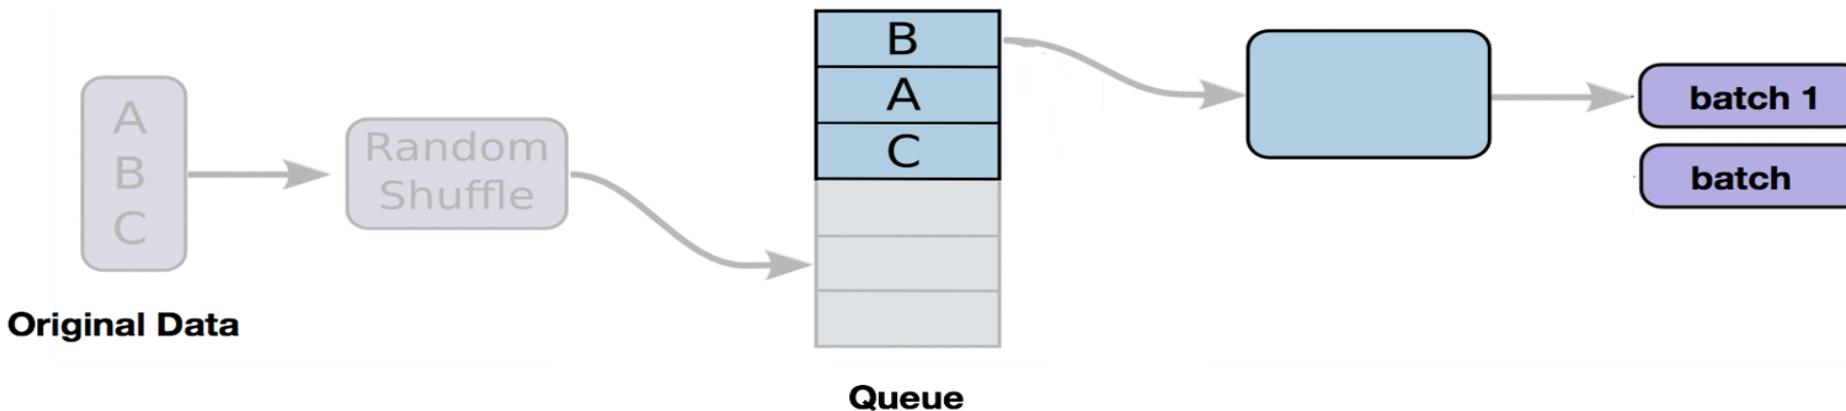
- one **epoch** = one forward pass and one backward pass of *all* the training examples
- **batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- number of **iterations** = number of passes, each pass using [batch size] number of examples.  
To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).

Example: if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch.



HAMIM®

# Data Loader



```
for i, data in enumerate(train_loader, 0):
    # get the inputs
    inputs, labels = data

    # wrap them in Variable
    inputs, labels = Variable(inputs),
    Variable(labels)

    # Run your training process
    print(epoch, i, "inputs", inputs.data,
          "labels", labels.data)
```

# Custom Data Loader

```
class DiabetesDataset(Dataset):
    """ Diabetes dataset."""

    # Initialize your data, download, etc.
    def __init__(self):
        1 download, read data, etc.

    def __getitem__(self, index):
        return
        2 return one item on the index

    def __len__(self):
        return

dataset = DiabetesDataset()
train_loader = DataLoader(dataset=dataset,
                         batch_size=32,
                         shuffle=True,
                         num_workers=2)
        3 return the data length
```



# Custom DataLoader

```
class DiabetesDataset(Dataset):
    """ Diabetes dataset."""

    # Initialize your data, download, etc.
    def __init__(self):
        xy = np.loadtxt('data-diabetes.csv', delimiter=',', dtype=np.float32)
        self.len = xy.shape[0]
        self.x_data = torch.from_numpy(xy[:, 0:-1])
        self.y_data = torch.from_numpy(xy[:, [-1]])

    def __getitem__(self, index):
        return self.x_data[index], self.y_data[index]

    def __len__(self):
        return self.len

dataset = DiabetesDataset()
train_loader = DataLoader(dataset=dataset,
                         batch_size=32,
                         shuffle=True,
                         num_workers=2)
```

# Using DataLoader

```
dataset = DiabetesDataset()
train_loader = DataLoader(dataset=dataset, batch_size=32, shuffle=True, num_workers=2)

# Training Loop
for epoch in range(2):
    for i, data in enumerate(train_loader, 0):
        # get the inputs
        inputs, labels = data

        # wrap them in Variable
        inputs, labels = Variable(inputs), Variable(labels)

        # Forward pass: Compute predicted y by passing x to the model
        y_pred = model(inputs)

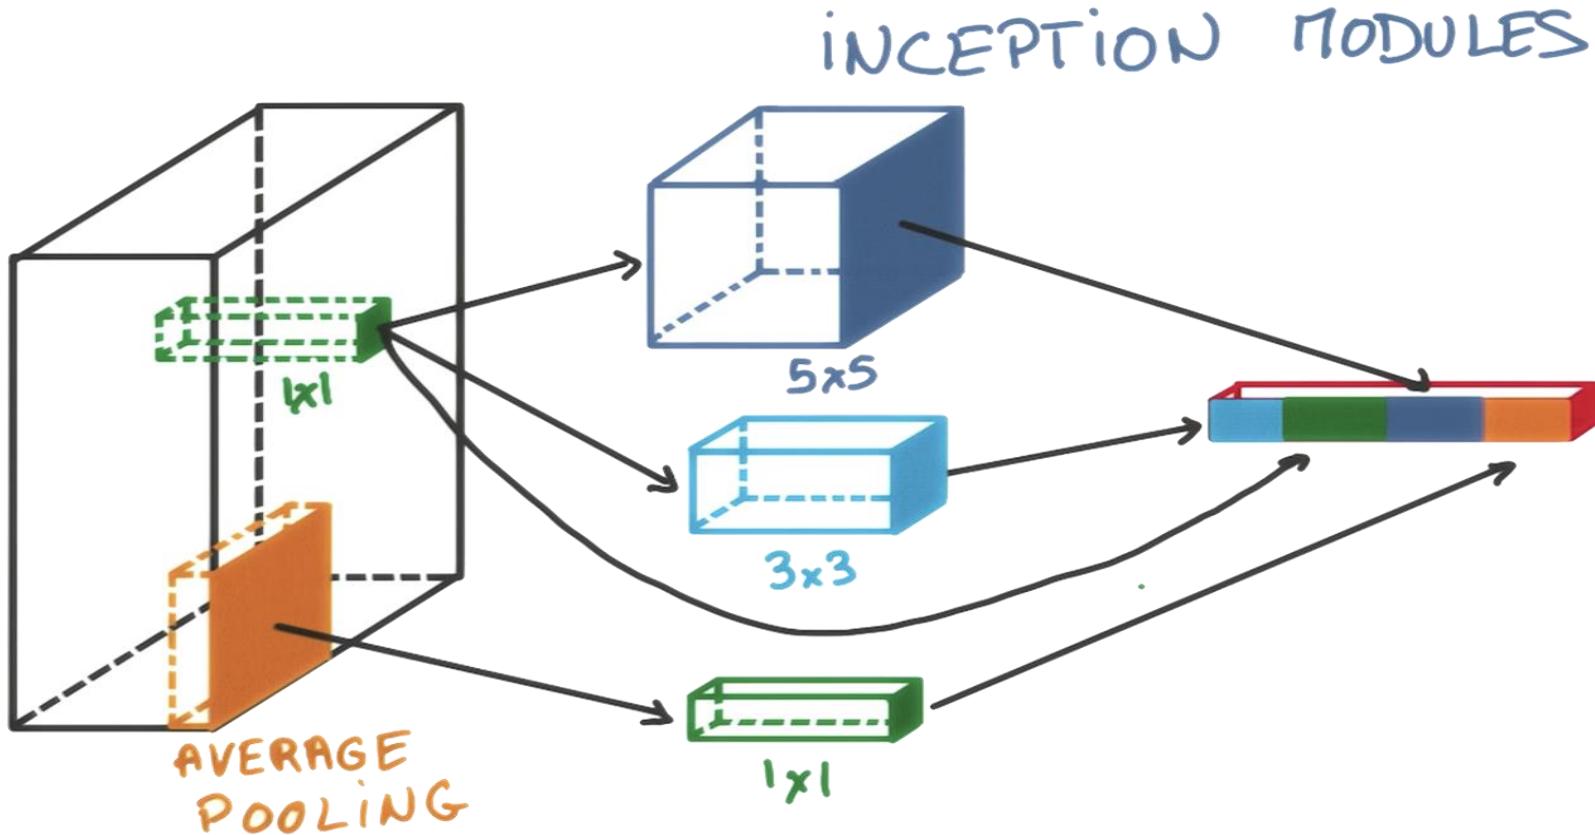
        # Compute and print loss
        loss = criterion(y_pred, labels)
        print(epoch, i, loss.data[0])

        # Zero gradients, perform a backward pass, and update the weights.
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```



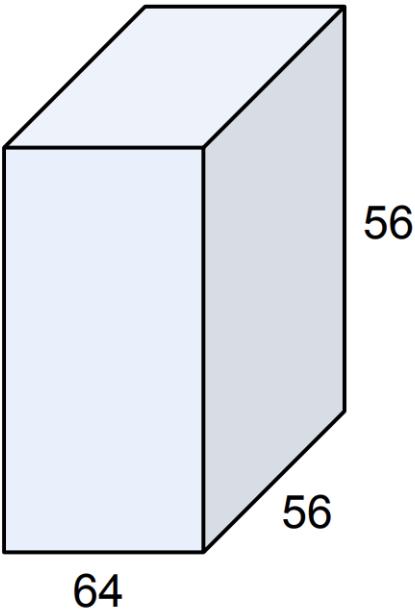
HAMIM®

# Inception Modules



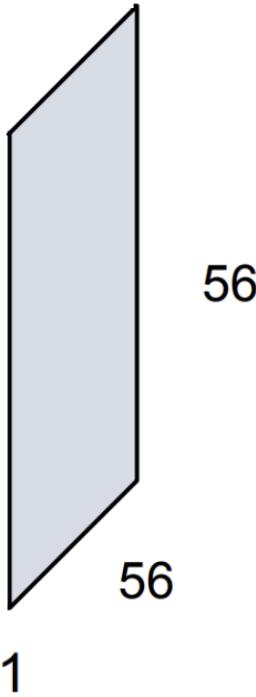
HAMIM®

# Why $1 \times 1$ convolution?



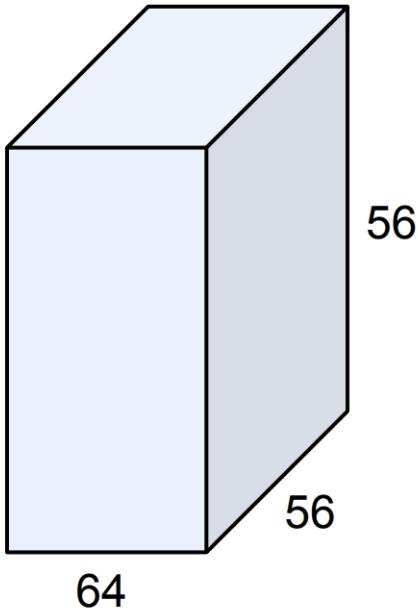
$1 \times 1$  CONV

(each filter has size  
 $1 \times 1 \times 64$ , and performs a  
64-dimensional dot  
product)

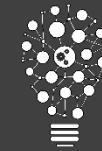
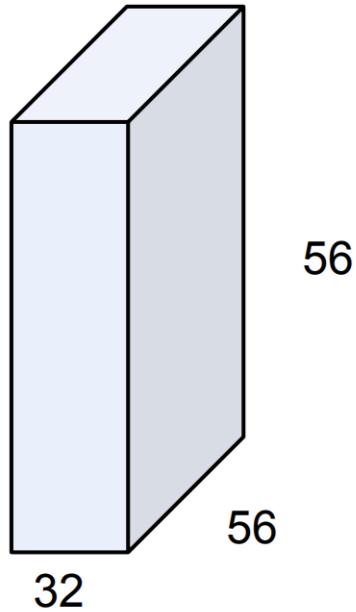


HAMIM®

# Why $1 \times 1$ convolution?



1x1 CONV  
with 32 filters  
—————  
(each filter has size  
 $1 \times 1 \times 64$ , and performs a  
64-dimensional dot product)



HAMIM®

# Why $1 \times 1$ convolution?

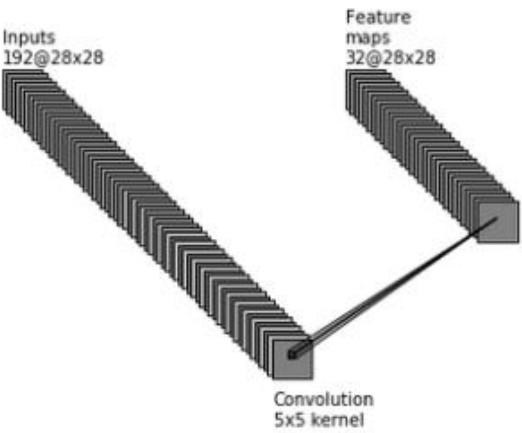


Figure 6.  $5 \times 5$  convolutions inside the Inception module using the naive model

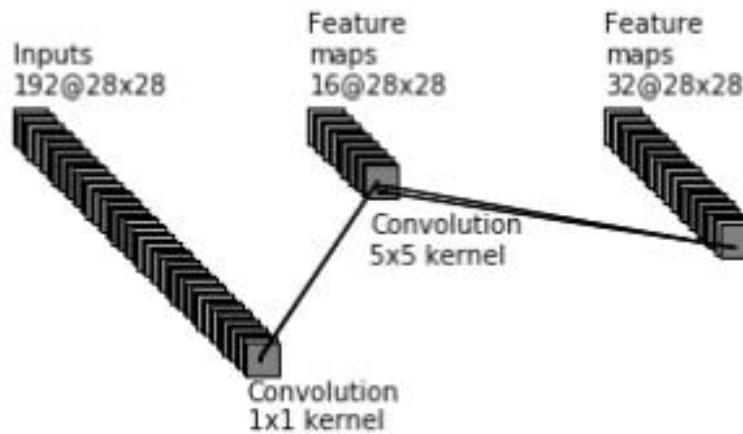


Figure 7.  $1 \times 1$  convolutions serve as the dimensionality reducers that limit the number of expensive  $5 \times 5$  convolutions that follow

# Why 1x1 convolution?

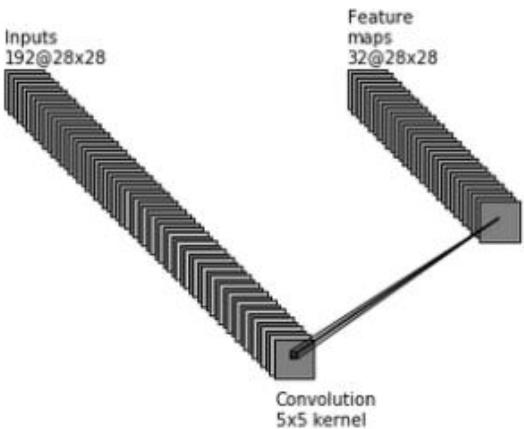


Figure 6.  $5 \times 5$  convolutions inside the Inception module using the naive model

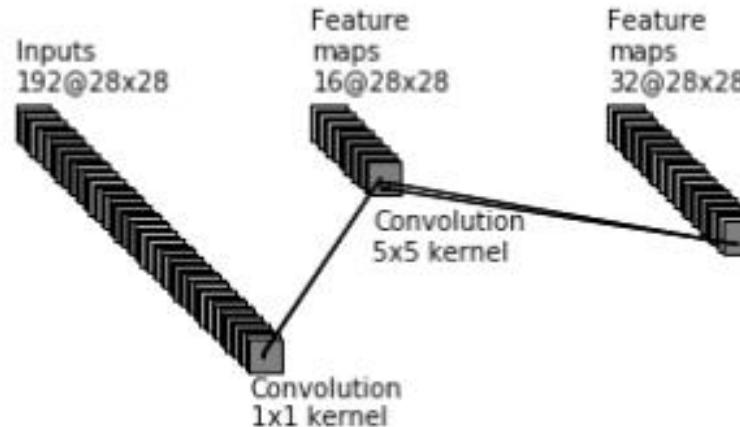


Figure 7.  $1 \times 1$  convolutions serve as the dimensionality reducers that limit the number of expensive  $5 \times 5$  convolutions that follow

Without 1x1 convolution:  
 $5^2 * 28^2 * 192 * 32 = 120,422,400$  operations

# Why 1x1 convolution?

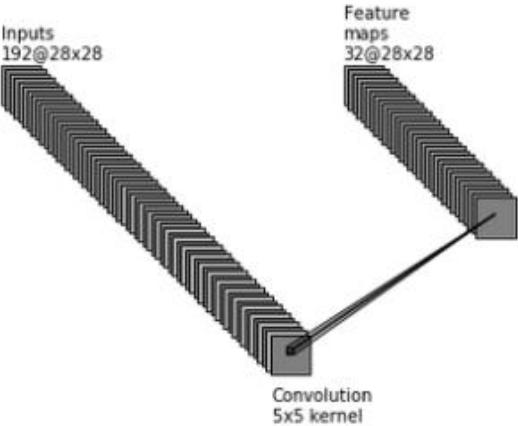


Figure 6. 5×5 convolutions inside the Inception module using the naive model

Without 1x1 convolution:  
 $5^2 * 28^2 * 192 * 32 = 120,422,400$  operations

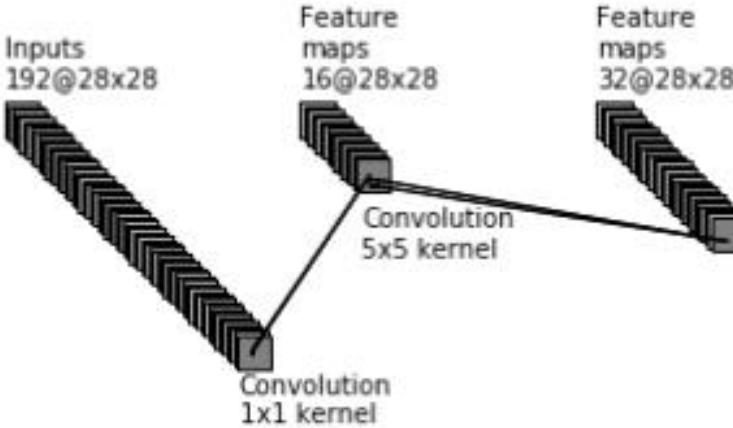
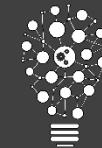
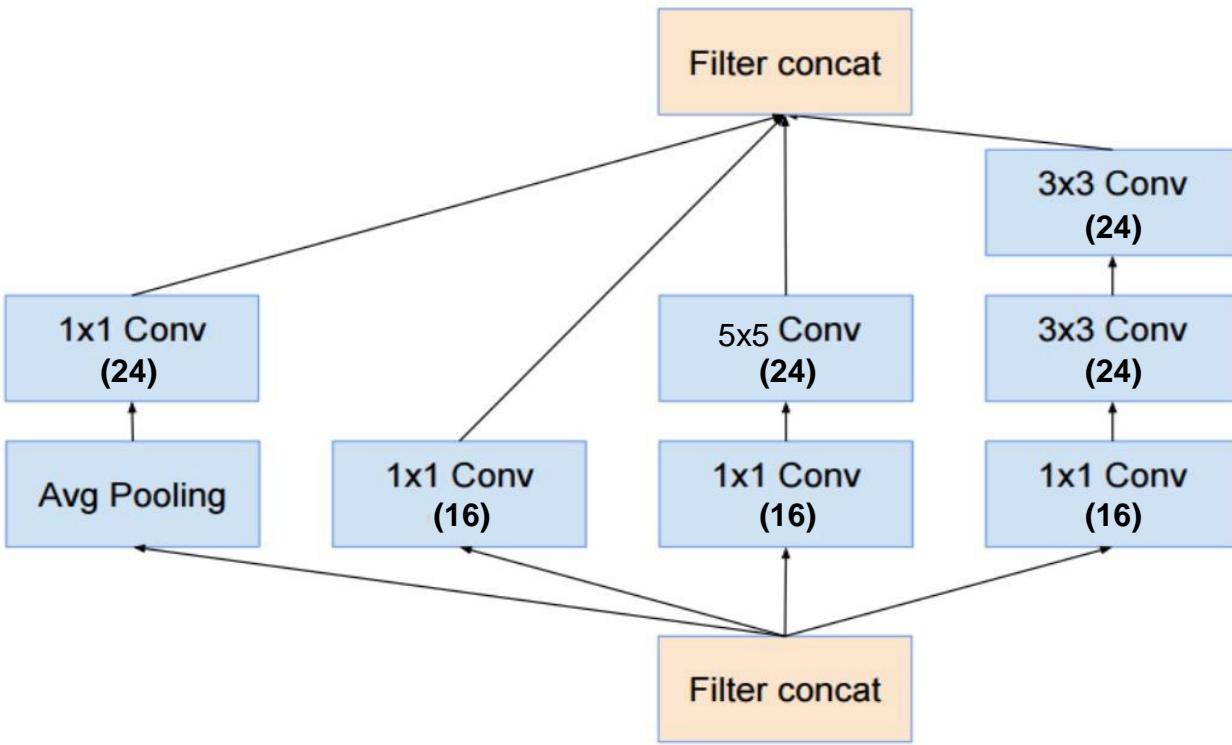


Figure 7. 1×1 convolutions serve as the dimensionality reducers that limit the number of expensive 5×5 convolutions that follow

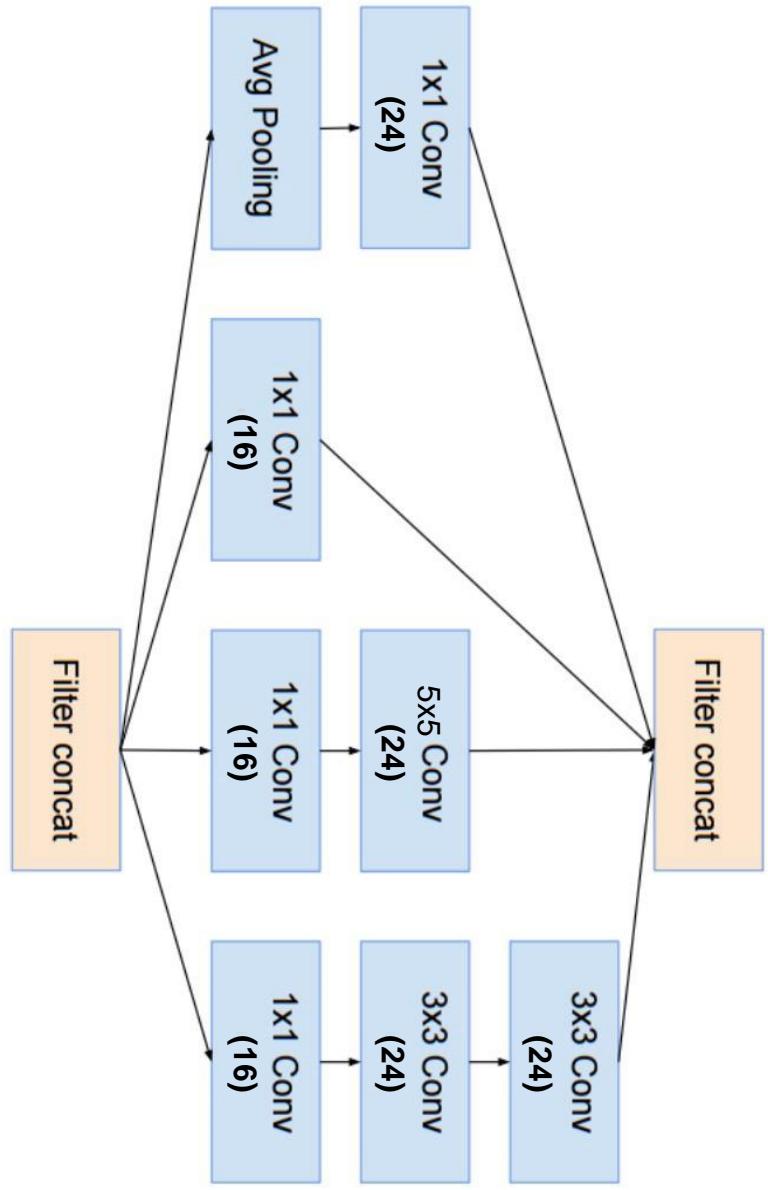
With 1x1 convolution:  
 $1^2 * 28^2 * 192 * 16$   
 $+ 5^2 * 28^2 * 16 * 32$   
 $= 12,443,648$  operations

<https://hacktilldawn.com/2016/09/25/inception-modules-explained-and-implemented/>

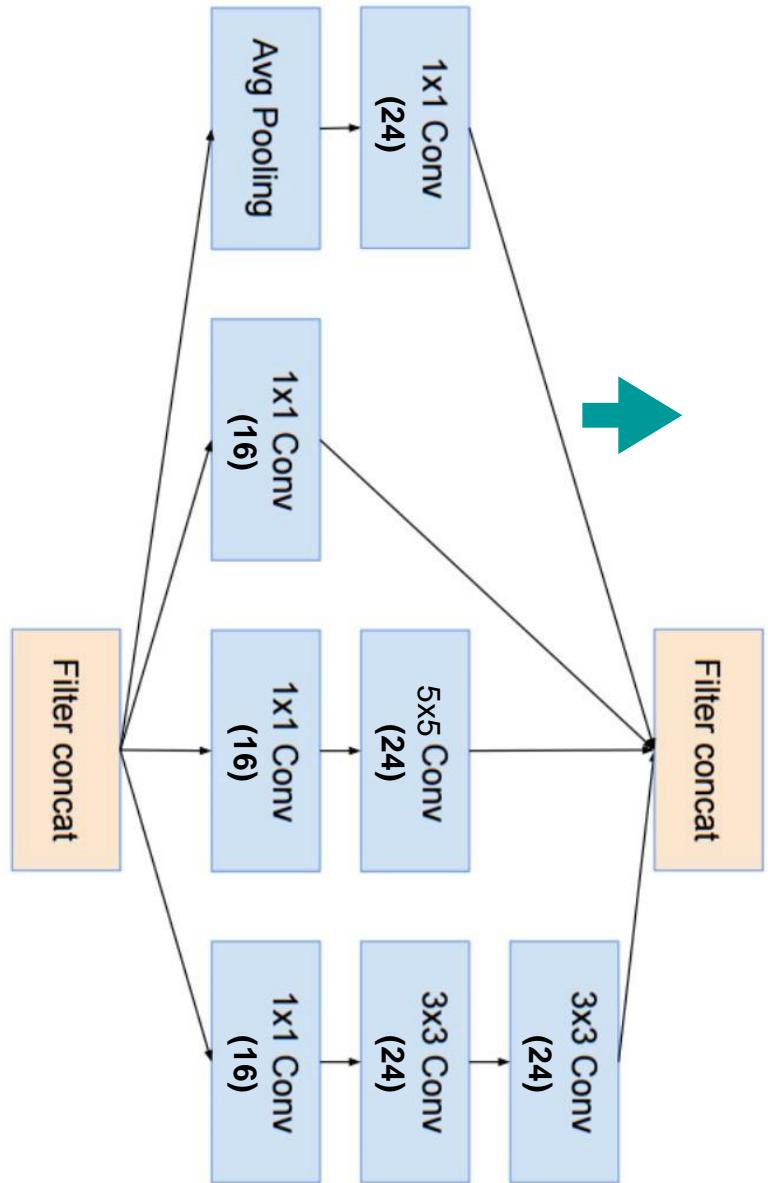
# Inception Module



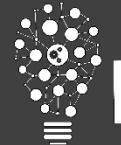
HAMIM®



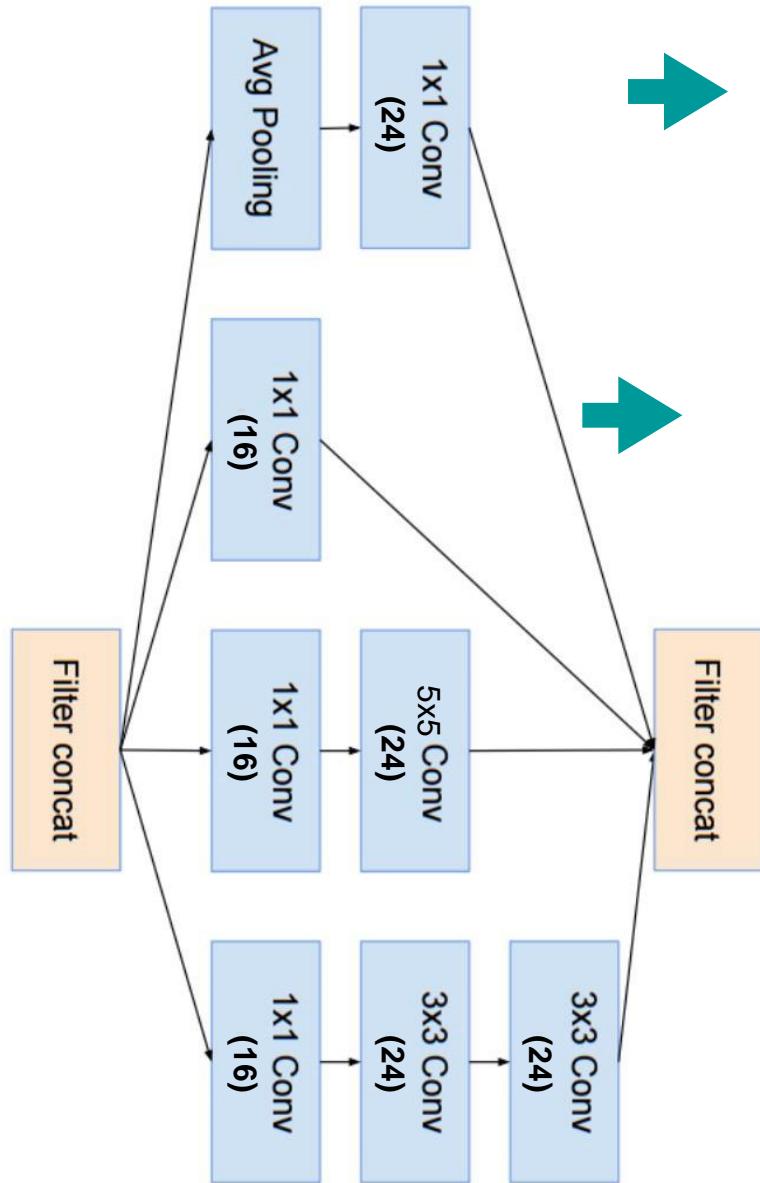
HAMIM®



```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)  
branch1x1 = self.branch1x1(x)
```



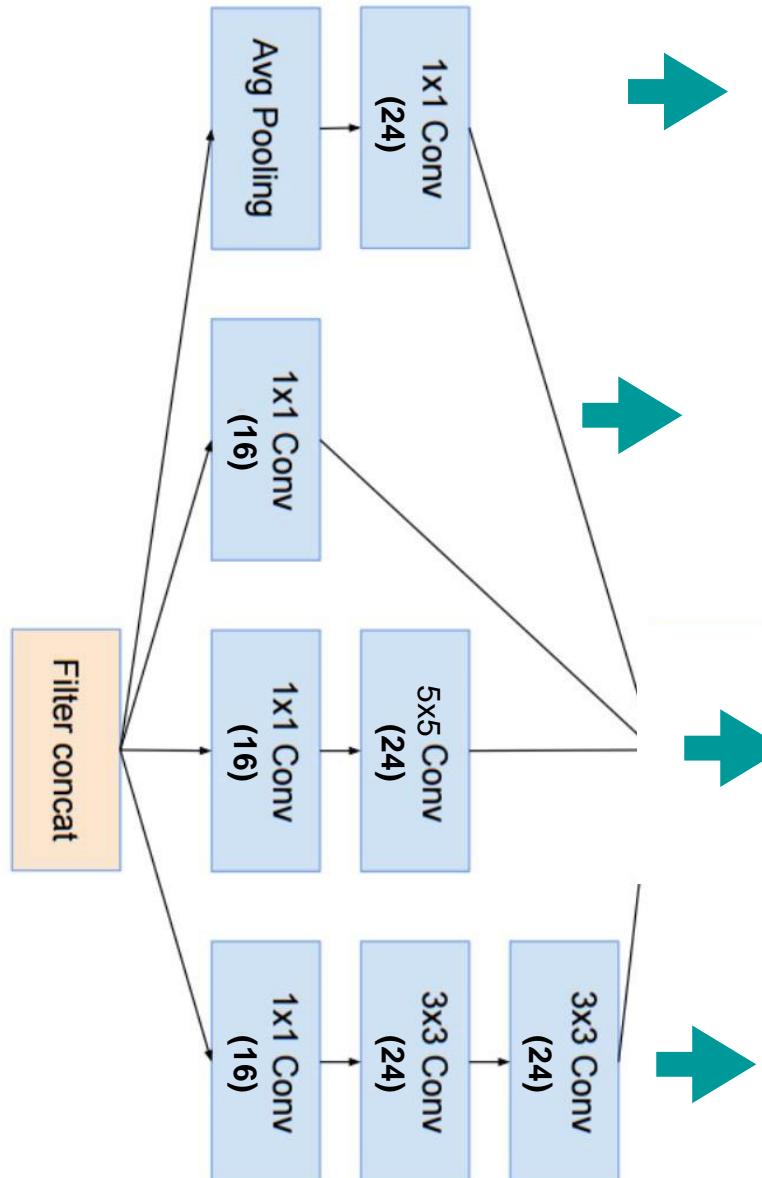
HAMIM®



```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)
```

```
branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)
branch1x1 = self.branch1x1(x)
```



```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)
```

```
branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)
branch1x1 = self.branch1x1(x)
```

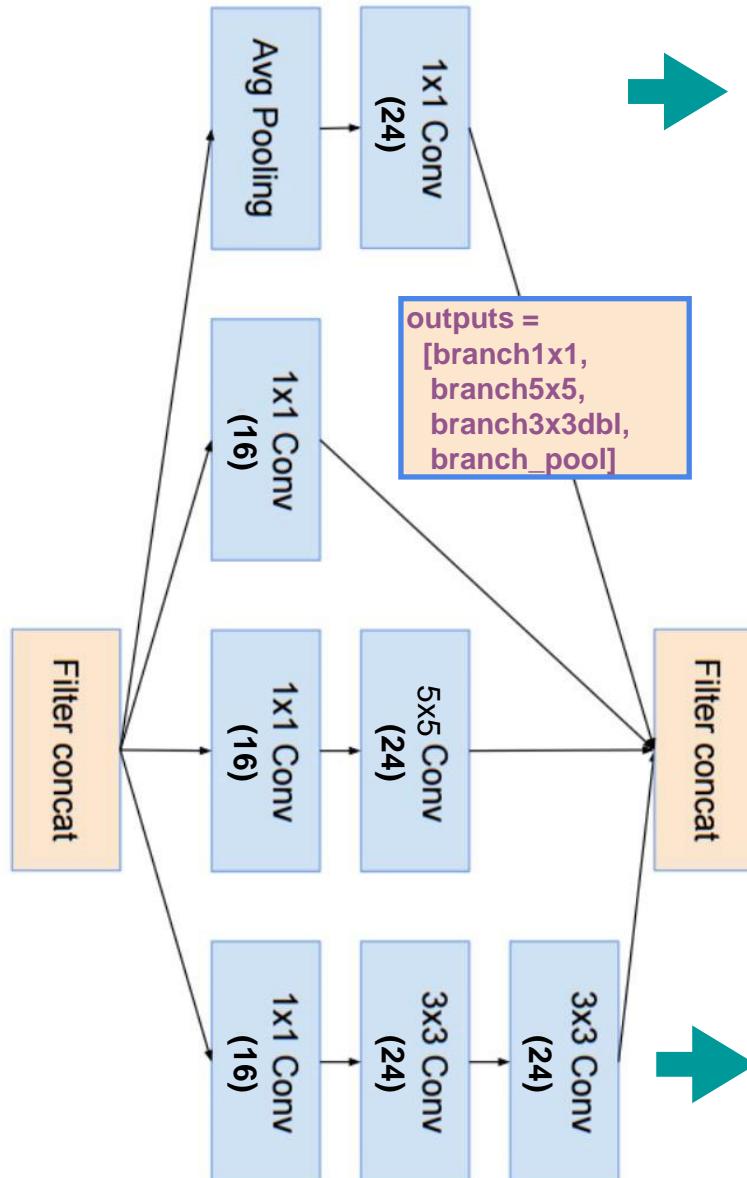
```
self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)
```

```
branch5x5 = self.branch5x5_1(x)
branch5x5 = self.branch5x5_2(branch5x5)
```

```
self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)
```

```
branch3x3dbl = self.branch3x3dbl_1(x)
branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)
```





```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)
```

```
branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```

```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
branch1x1 = self.branch1x1(x)
```

```
self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)
```

```
branch5x5 = self.branch5x5_1(x)
```

```
branch5x5 = self.branch5x5_2(branch5x5)
```

```
self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
```

```
self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
```

```
self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)
```

```
branch3x3dbl = self.branch3x3dbl_1(x)
```

```
branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
```

```
branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)
```



# Inception Module

```
class InceptionA(nn.Module):
    def __init__(self, in_channels):
        super(InceptionA, self).__init__()
        self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

        self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)

        self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
        self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)

        self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

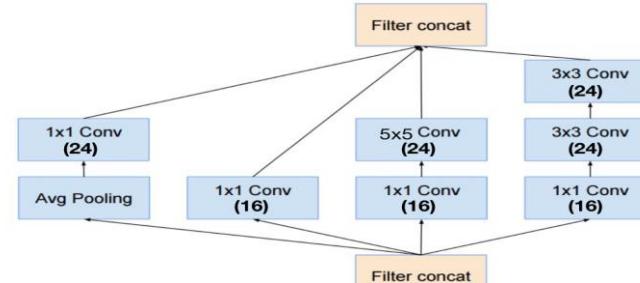
    def forward(self, x):
        branch1x1 = self.branch1x1(x)

        branch5x5 = self.branch5x5_1(x)
        branch5x5 = self.branch5x5_2(branch5x5)

        branch3x3dbl = self.branch3x3dbl_1(x)
        branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
        branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)

        branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
        branch_pool = self.branch_pool(branch_pool)

        outputs = [branch1x1, branch5x5, branch3x3dbl, branch_pool]
        return torch.cat(outputs, 1)
```



```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(88, 20, kernel_size=5)

        self.incept1 = InceptionA(in_channels=10)
        self.incept2 = InceptionA(in_channels=20)

        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(1408, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = self.incept1(x)
        x = F.relu(self.mp(self.conv2(x)))
        x = self.incept2(x)
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

# Deeper ???

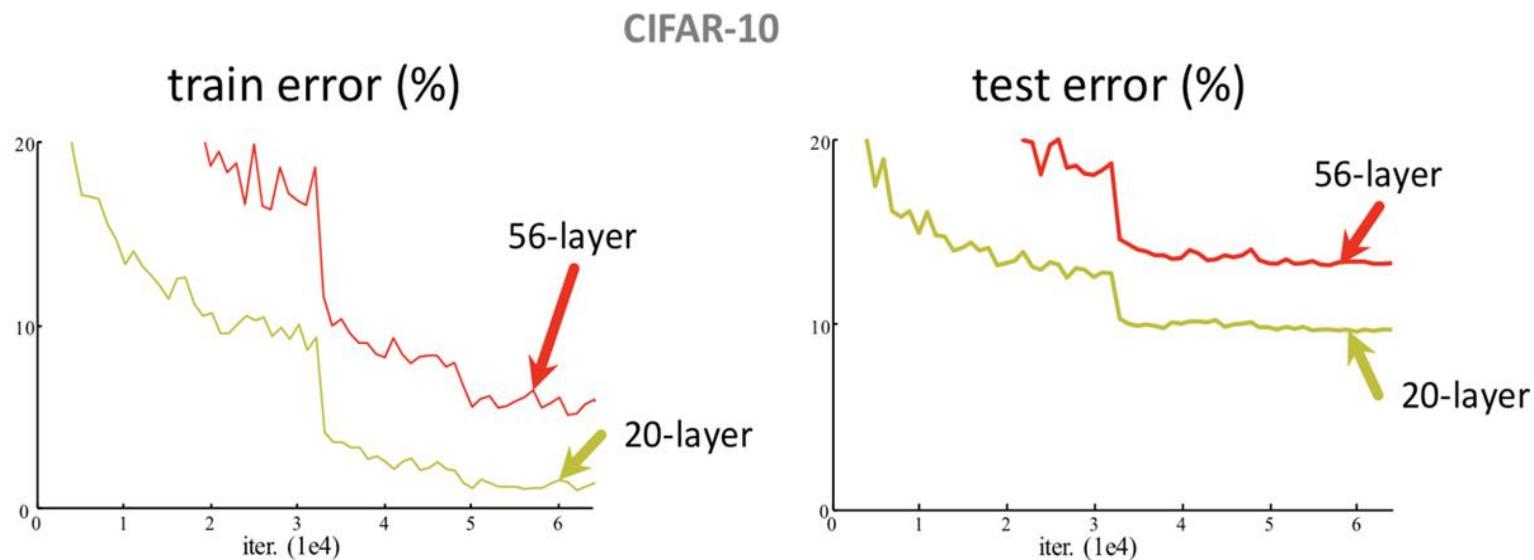


Well said Leo, well said



HAMIM®

# Can we just go deeper, keep stacking layers?



- Plain nets: stacking 3x3 conv layers...
- 56-layer net has **higher training error** and test error than 20-layer net

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". CVPR 2016.